Deep Learning for Computer Vision

Explaining CNNs: Recent Methods

Vineeth N Balasubramanian

Department of Computer Science and Engineering
Indian Institute of Technology, Hyderabad

Recall: Explaining using Vanilla Gradients[1]

- Forward pass the data $\mathbf{x}$, to get $y = f(\mathbf{x})$, where y is DNN's output corresponding to a given class.

[1] Simonyan et al, Deep Inside Convolutional Networks:Visualising Image Classification Models and Saliency Maps, ICLRW 2014

Recall: Explaining using Vanilla Gradients[1]

- Forward pass the data $\mathbf{x}$, to get $y = f(\mathbf{x})$, where y is DNN's output corresponding to a given class.
- Backward pass to input layer to get the gradient $\frac{\partial y}{\partial \mathbf{x}}$.

[1] Simonyan et al, Deep Inside Convolutional Networks:Visualising Image Classification Models and Saliency Maps, ICLRW 2014

Recall: Explaining using Vanilla Gradients[1]

- Forward pass the data $\mathbf{x}$, to get $y = f(\mathbf{x})$, where y is DNN's output corresponding to a given class.
- Backward pass to input layer to get the gradient $\frac{\partial y}{\partial \mathbf{x}}$.

Original image *(left)*; Vanilla Gradients Attribution map *(right)*

Is this enough to explain a Deep Neural Network?

[1] Simonyan et al, Deep Inside Convolutional Networks:Visualising Image Classification Models and Saliency Maps, ICLRW 2014

Recall: Explaining using Vanilla Gradients[1]

- Forward pass the data $\mathbf{x}$, to get $y = f(\mathbf{x})$, where y is DNN's output corresponding to a given class.
- Backward pass to input layer to get the gradient $\frac{\partial y}{\partial \mathbf{x}}$.

Original image *(left)*; Vanilla Gradients Attribution map *(right)*

Is this enough to explain a Deep Neural Network?
Not always!

[1] Simonyan et al, Deep Inside Convolutional Networks:Visualising Image Classification Models and Saliency Maps, ICLRW 2014

Saturation Problem!

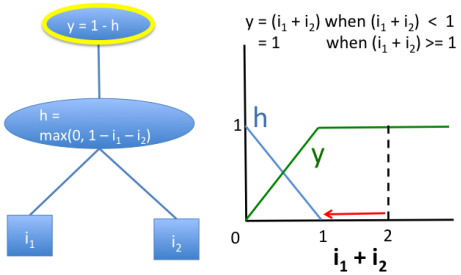

Illustration of saturation problem

Saturation Problem!

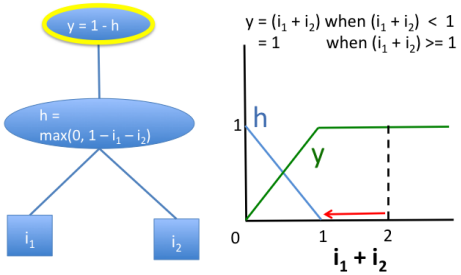

Illustration of saturation problem

- Gradient of h w.r.t both i_1 and i_2 is zero when $i_1 + i_2 > 1$ (causing both gradients and Guided Backprop to be zero)

Saturation Problem!

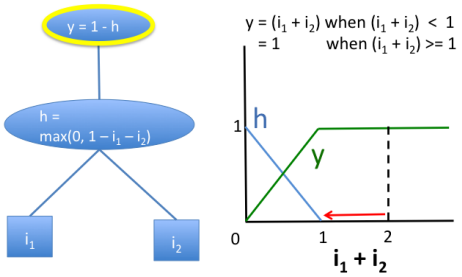

Illustration of saturation problem

- Gradient of h w.r.t both i_1 and i_2 is zero when $i_1 + i_2 > 1$ (causing both gradients and Guided Backprop to be zero)
- Gradient of y w.r.t. h is negative (causing Guided Backprop and deconvolutional networks to assign zero importance)

Deep Lift[2]

- **Idea**: Instead of gradients, measure difference in output from some 'reference' output (Δt) in terms of difference of input from some 'reference' input (Δx_i).

[2]Shrikumar et al, Learning important features through propagating activation differences, ICML 2017

Deep Lift[2]

- **Idea**: Instead of gradients, measure difference in output from some 'reference' output (Δt) in terms of difference of input from some 'reference' input (Δx_i).
- Assigns contribution scores $C_{\Delta x_i \Delta t}$ s.t. $\sum_{i=1}^{n} C_{\Delta x_i \Delta t} = \Delta t$

[2]Shrikumar et al, Learning important features through propagating activation differences, ICML 2017

Deep Lift[2]

- **Idea**: Instead of gradients, measure difference in output from some 'reference' output (Δt) in terms of difference of input from some 'reference' input (Δx_i).
- Assigns contribution scores $C_{\Delta x_i \Delta t}$ s.t. $\sum_{i=1}^{n} C_{\Delta x_i \Delta t} = \Delta t$

[2]Shrikumar et al, Learning important features through propagating activation differences, ICML 2017

Deep Lift[2]

- **Idea**: Instead of gradients, measure difference in output from some 'reference' output (Δt) in terms of difference of input from some 'reference' input (Δx_i).
- Assigns contribution scores $C_{\Delta x_i \Delta t}$ s.t. $\sum_{i=1}^{n} C_{\Delta x_i \Delta t} = \Delta t$

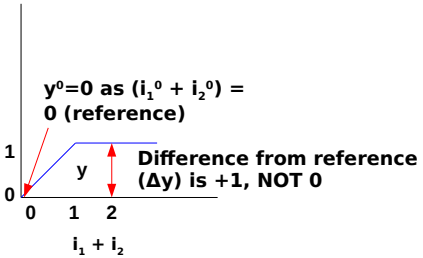

DeepLift overcomes saturation problem

[2]Shrikumar et al, Learning important features through propagating activation differences, ICML 2017

Deep Lift: Rescale Rule

- **Idea**: Start from output layer L & proceed backwards layer by layer, redistributing the difference of prediction score from baseline, until input layer is reached

Deep Lift: Rescale Rule

- **Idea**: Start from output layer L & proceed backwards layer by layer, redistributing the difference of prediction score from baseline, until input layer is reached
- **Notations**:
  - $z_{ji} = w_{ji}^{(l+1,l)} x_i^l$: weighted activation of a neuron i onto neuron j in the next layer

Deep Lift: Rescale Rule

- **Idea**: Start from output layer L & proceed backwards layer by layer, redistributing the difference of prediction score from baseline, until input layer is reached
- **Notations**:
  - $z_{ji} = w_{ji}^{(l+1,l)} x_i^l$: weighted activation of a neuron i onto neuron j in the next layer
  - $\overline{z_{ji}} = w_{ji}^{(l+1,l)} \bar{x}_i^l$: weighted activation of a neuron i onto neuron j in the next layer, when baseline $\overline{x}$ fed as the input.

Deep Lift: Rescale Rule

- **Idea**: Start from output layer L & proceed backwards layer by layer, redistributing the difference of prediction score from baseline, until input layer is reached
- **Notations**:
  - $z_{ji} = w_{ji}^{(l+1,l)} x_i^l$: weighted activation of a neuron i onto neuron j in the next layer
  - $\overline{z_{ji}} = w_{ji}^{(l+1,l)} \bar{x}_i^l$: weighted activation of a neuron i onto neuron j in the next layer, when baseline $\overline{x}$ fed as the input.
  - $r_i^{(l)}$: relevance of unit i of layer l.

Deep Lift: Rescale Rule

- **Idea**: Start from output layer L & proceed backwards layer by layer, redistributing the difference of prediction score from baseline, until input layer is reached
- **Notations**:
  - $z_{ji} = w_{ji}^{(l+1,l)} x_i^l$: weighted activation of a neuron i onto neuron j in the next layer
  - $\overline{z_{ji}} = w_{ji}^{(l+1,l)} \bar{x}_i^l$: weighted activation of a neuron i onto neuron j in the next layer, when baseline $\overline{x}$ fed as the input.
  - $r_i^{(l)}$: relevance of unit i of layer l.
- $r_i^{(L)}$: $\begin{cases} = \mathrm{y}_i(x) - y_i(\bar{x}) \text{ if unit } i \text{ is target unit of interest} \\ = 0 \text{ otherwise} \end{cases}$

Deep Lift: Rescale Rule

- **Idea**: Start from output layer L & proceed backwards layer by layer, redistributing the difference of prediction score from baseline, until input layer is reached
- **Notations**:
  - $z_{ji} = w_{ji}^{(l+1,l)} x_i^l$: weighted activation of a neuron i onto neuron j in the next layer
  - $\overline{z_{ji}} = w_{ji}^{(l+1,l)} \bar{x}_i^l$: weighted activation of a neuron i onto neuron j in the next layer, when baseline $\overline{x}$ fed as the input.
  - $r_i^{(l)}$: relevance of unit i of layer l.
- $r_i^{(L)}$: $\begin{cases} = \mathrm{y}_i(x) - y_i(\bar{x}) \text{ if unit } i \text{ is target unit of interest} \\ = 0 \text{ otherwise} \end{cases}$
- $r_i^{(l)} = \sum_j \frac{z_{ji} - \overline{z_{ji}}}{\sum_{i'}(z_{ji'} - \overline{z_{ji'}})} r_j^{(l+1)}$

IG: Integrated Gradients[3]


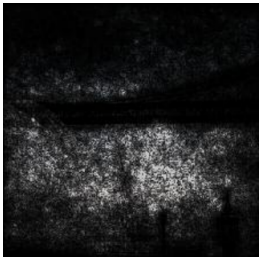

Image of Fireboat *(left)*, Vanilla Gradients *(right)*

[3]Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017

IG: Integrated Gradients[3]

Image of Fireboat *(left)*, Vanilla Gradients *(right)*

- Due to saturation problem, vanilla gradients highlight regions irrelevant to fireboat

[3]Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017

IG: Integrated Gradients[3]

Image of Fireboat *(left)*, Vanilla Gradients *(right)*

- Due to saturation problem, vanilla gradients highlight regions irrelevant to fireboat
- IG overcomes problem of saturating gradients by cumulating gradients at different pixel intensities, α's.

[3]Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017

IG: Integrated Gradients

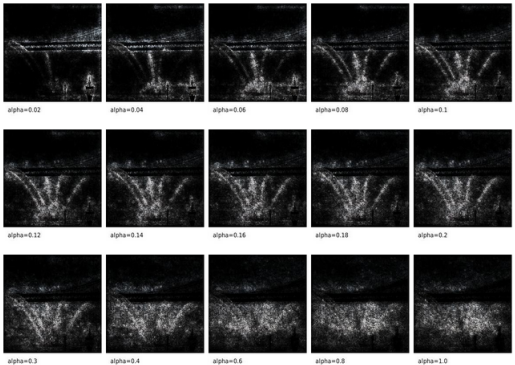

Gradients at increasing α values from top-left to bottom-right

IG: Integrated Gradients

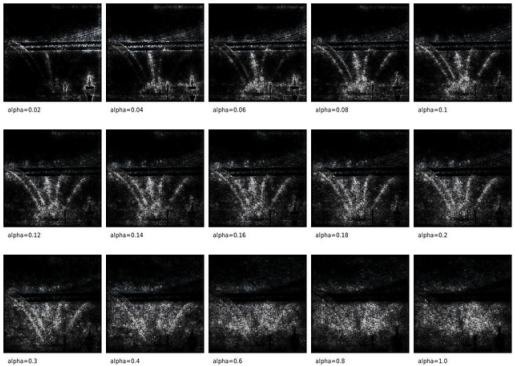

Gradients at increasing α values from top-left to bottom-right

- Region of importance is changing with increasing α. To get a more realistic picture of what is going on, cumulate these gradients using **path integral**

IG: Integrated Gradients

- **Integrated gradient** along i^{th} dimension for input x and baseline x' given by:

$$\mathsf{IG}_i(x) ::= (x_i - x'_i) \int_{\alpha=0}^{1} \frac{\partial f(x' + \alpha(x - x'))}{\partial x_i} \partial\alpha$$

IG: Integrated Gradients

- **Integrated gradient** along i^{th} dimension for input x and baseline x' given by:

$$\mathsf{IG}_i(x) ::= (x_i - x'_i) \int_{\alpha=0}^{1} \frac{\partial f(x' + \alpha(x - x'))}{\partial x_i} \partial\alpha$$

- $\mathsf{IG}_i^{\mathsf{approx}}(x) ::= (x_i - x'_i) \sum_{k=1}^{m} \frac{\partial f(x' + \frac{k}{m}(x - x'))}{\partial x_i} \frac{1}{m}$ where m is a hyperparameter.

IG: Integrated Gradients

- **Integrated gradient** along i^{th} dimension for input x and baseline x' given by:

$$\mathsf{IG}_i(x) ::= (x_i - x_i') \int_{\alpha=0}^{1} \frac{\partial f(x' + \alpha(x - x'))}{\partial x_i} \partial\alpha$$

- $\mathsf{IG}_i^{\mathsf{approx}}(x) ::= (x_i - x_i') \sum_{k=1}^{m} \frac{\partial f(x' + \frac{k}{m}(x - x'))}{\partial x_i} \frac{1}{m}$ where m is a hyperparameter.

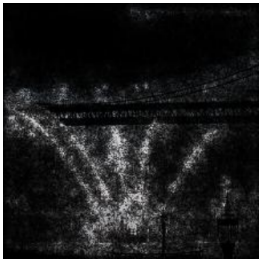

IG attribution map

SmoothGrad[4]

- Add pixel-wise Gaussian noise to many copies of the image, and average resulting gradients.

[4] Smilkov et al, SmoothGrad: removing noise by adding noise, ICMLW 2017

SmoothGrad[4]

- Add pixel-wise Gaussian noise to many copies of the image, and average resulting gradients.
- Removes noise from saliency map by adding noise!

[4] Smilkov et al, SmoothGrad: removing noise by adding noise, ICMLW 2017

SmoothGrad[4]

- Add pixel-wise Gaussian noise to many copies of the image, and average resulting gradients.
- Removes noise from saliency map by adding noise!
- Besides vanilla gradients, other attribution methods also have their SmoothGrad counterparts, e.g. Smooth Integrated Gradients

[4] Smilkov et al, SmoothGrad: removing noise by adding noise, ICMLW 2017

SmoothGrad[4]

- Add pixel-wise Gaussian noise to many copies of the image, and average resulting gradients.
- Removes noise from saliency map by adding noise!
- Besides vanilla gradients, other attribution methods also have their SmoothGrad counterparts, e.g. Smooth Integrated Gradients

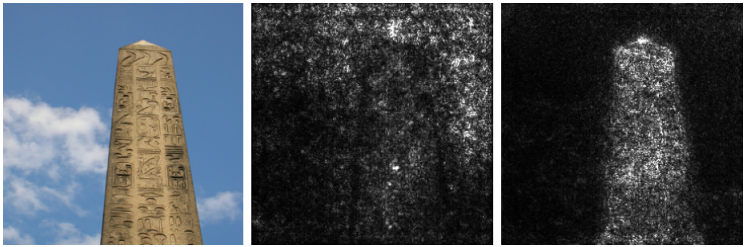

Original Image *(left)*, Vanilla Gradients *(center)*, SmoothGrad *(right)*

[4] Smilkov et al, SmoothGrad: removing noise by adding noise, ICMLW 2017

Recent Variant of IG: XRAI[5]

- Get attribution map given by IG

[5] Kapishnikov et al, XRAI: Better Attributions Through Regions, ICCV 2019

Recent Variant of IG: XRAI[5]

- Get attribution map given by IG
- Over-segment the image

[5] Kapishnikov et al, XRAI: Better Attributions Through Regions, ICCV 2019

Recent Variant of IG: XRAI[5]

- Get attribution map given by IG
- Over-segment the image
- Start with an empty mask

[5] Kapishnikov et al, XRAI: Better Attributions Through Regions, ICCV 2019

Recent Variant of IG: XRAI[5]

- Get attribution map given by IG
- Over-segment the image
- Start with an empty mask
- Populate this mask by selectively adding segments that yield maximum gain in total attributions per area

[5] Kapishnikov et al, XRAI: Better Attributions Through Regions, ICCV 2019

Recent Variant of IG: XRAI[5]

- Get attribution map given by IG
- Over-segment the image
- Start with an empty mask
- Populate this mask by selectively adding segments that yield maximum gain in total attributions per area

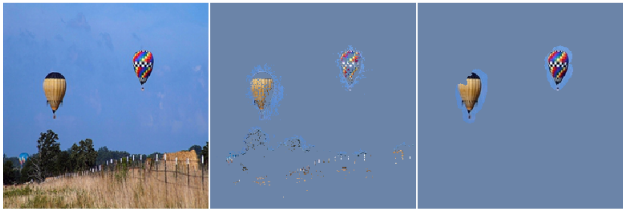

Original image *(left)*, IG *(center)*, XRAI *(right)*

[5]Kapishnikov et al, XRAI: Better Attributions Through Regions, ICCV 2019

Recent Variant of IG: XRAI[6]

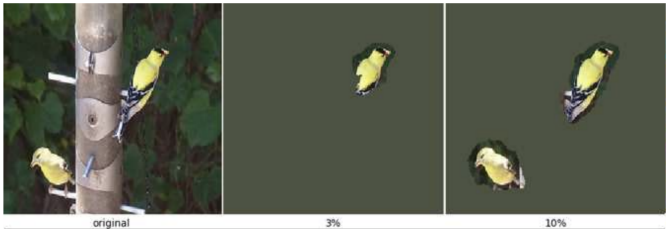

[6]Kapishnikov et al, XRAI: Better Attributions Through Regions, ICCV 2019

LIME: Local Interpretable Model-agnostic Explanations[7]

- **Idea**: Approximate underlying model locally by an interpretable (typically linear) one

[7] Ribiero et al, Why Should I Trust You?: Explaining the Predictions of Any Classifier, KDD 2016

LIME: Local Interpretable Model-agnostic Explanations[7]

- **Idea**: Approximate underlying model locally by an interpretable (typically linear) one
- Interpretable models are trained on small perturbations of original instance

[7] Ribiero et al, Why Should I Trust You?: Explaining the Predictions of Any Classifier, KDD 2016

LIME: Local Interpretable Model-agnostic Explanations[7]

- **Idea**: Approximate underlying model locally by an interpretable (typically linear) one
- Interpretable models are trained on small perturbations of original instance

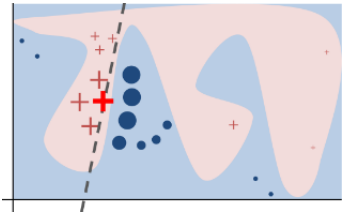

Intuition for LIME

- *Blue/Pink background:* black box model's decision function f

[7] Ribiero et al, Why Should I Trust You?: Explaining the Predictions of Any Classifier, KDD 2016

LIME: Local Interpretable Model-agnostic Explanations[7]

- **Idea**: Approximate underlying model locally by an interpretable (typically linear) one
- Interpretable models are trained on small perturbations of original instance

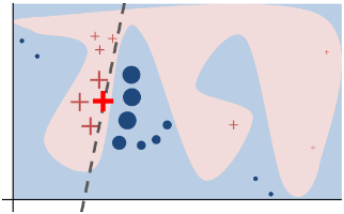

Intuition for LIME

- *Blue/Pink background:* black box model's decision function f
- *Bold red cross:* instance being explained

[7] Ribiero et al, Why Should I Trust You?: Explaining the Predictions of Any Classifier, KDD 2016

LIME: Local Interpretable Model-agnostic Explanations[7]

- **Idea**: Approximate underlying model locally by an interpretable (typically linear) one
- Interpretable models are trained on small perturbations of original instance

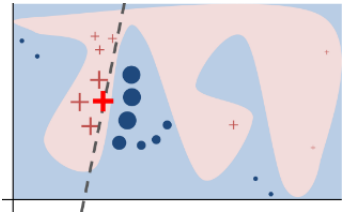

Intuition for LIME

- *Blue/Pink background:* black box model's decision function f
- *Bold red cross:* instance being explained
- *Dashed line:* learned explanation

[7] Ribiero et al, Why Should I Trust You?: Explaining the Predictions of Any Classifier, KDD 2016

LIME: Local Interpretable Model-agnostic Explanations

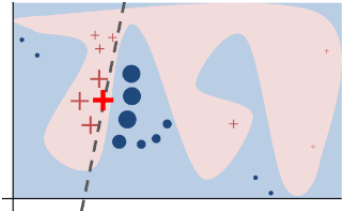

- Given a point x, let $z' \in \mathcal{Z}$ be a point obtained by perturbing one dimension (or region) in x; $\mathcal{Z}$ is the set of perturbations

LIME: Local Interpretable Model-agnostic Explanations

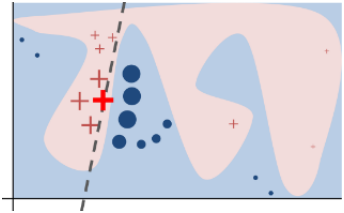

- Given a point x, let $z' \in \mathcal{Z}$ be a point obtained by perturbing one dimension (or region) in x; $\mathcal{Z}$ is the set of perturbations
- $\pi_x(z')$: proximity measure between instances z' and x, e.g. $\pi_x(z') = \exp(-D(x, z')^2/\sigma^2)$

LIME: Local Interpretable Model-agnostic Explanations

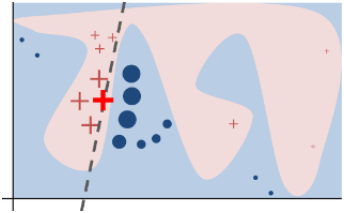

- Given a point x, let $z' \in \mathcal{Z}$ be a point obtained by perturbing one dimension (or region) in x; $\mathcal{Z}$ is the set of perturbations
- $\pi_x(z')$: proximity measure between instances z' and x, e.g. $\pi_x(z') = \exp(-D(x, z')^2/\sigma^2)$
- $f : \mathbb{R}^d \rightarrow \mathbb{R}$: model being explained

LIME: Local Interpretable Model-agnostic Explanations

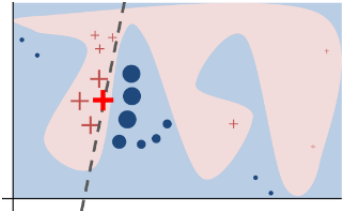

- Given a point x, let $z' \in \mathcal{Z}$ be a point obtained by perturbing one dimension (or region) in x; $\mathcal{Z}$ is the set of perturbations
- $\pi_x(z')$: proximity measure between instances z' and x, e.g. $\pi_x(z') = \exp(-D(x, z')^2/\sigma^2)$
- $f : \mathbb{R}^d \rightarrow \mathbb{R}$: model being explained
- Build a sparse linear model, $g(z') = w_g \cdot z'$

LIME: Local Interpretable Model-agnostic Explanations

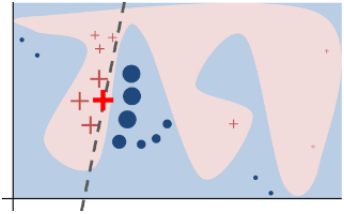

- Given a point x, let $z' \in \mathcal{Z}$ be a point obtained by perturbing one dimension (or region) in x; $\mathcal{Z}$ is the set of perturbations
- $\pi_x(z')$: proximity measure between instances z' and x, e.g. $\pi_x(z') = \exp(-D(x, z')^2/\sigma^2)$
- $f : \mathbb{R}^d \rightarrow \mathbb{R}$: model being explained
- Build a sparse linear model, $g(z') = w_g \cdot z'$
- Learn w_g to minimize:

$$\mathcal{L}(f, g, \pi_x) = \sum_{z, z' \in \mathcal{Z}} \pi_x(z)(f(z) - g(z'))^2$$

LIME: Local Interpretable Model-agnostic Explanations

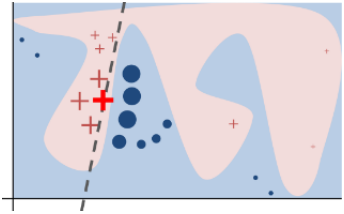

- Given a point x, let $z' \in \mathcal{Z}$ be a point obtained by perturbing one dimension (or region) in x; $\mathcal{Z}$ is the set of perturbations
- $\pi_x(z')$: proximity measure between instances z' and x, e.g. $\pi_x(z') = \exp(-D(x, z')^2/\sigma^2)$
- $f : \mathbb{R}^d \rightarrow \mathbb{R}$: model being explained
- Build a sparse linear model, $g(z') = w_g \cdot z'$
- Learn w_g to minimize:

$$\mathcal{L}(f, g, \pi_x) = \sum_{z, z' \in \mathcal{Z}} \pi_x(z)(f(z) - g(z'))^2$$

LIME: Local Interpretable Model-agnostic Explanations

(a) Original Image

(b) Explaining *Electric guitar*

(c) Explaining *Acoustic guitar*

(d) Explaining *Labrador*

LIME: Fidelity-Interpretability Trade-off

Notations:

- $\mathcal{L}(f, g, \pi_x)$: measure of how unfaithful g is in approximating f in the locality defined by π_x

LIME: Fidelity-Interpretability Trade-off

Notations:

- $\mathcal{L}(f, g, \pi_x)$: measure of how unfaithful g is in approximating f in the locality defined by π_x
- $g \in G$: model belonging to class of interpretable models, e.g. linear model $g(z') = w_g.z'$

LIME: Fidelity-Interpretability Trade-off

Notations:

- $\mathcal{L}(f, g, \pi_x)$: measure of how unfaithful g is in approximating f in the locality defined by π_x
- $g \in G$: model belonging to class of interpretable models, e.g. linear model $g(z') = w_g.z'$
- $\Omega(g)$: Complexity of g model

LIME: Fidelity-Interpretability Trade-off

Notations:

- $\mathcal{L}(f, g, \pi_x)$: measure of how unfaithful g is in approximating f in the locality defined by π_x
- $g \in G$: model belonging to class of interpretable models, e.g. linear model $g(z') = w_g.z'$
- $\Omega(g)$: Complexity of g model
  - Depth of trees in decision trees

LIME: Fidelity-Interpretability Trade-off

Notations:

- $\mathcal{L}(f, g, \pi_x)$: measure of how unfaithful g is in approximating f in the locality defined by π_x
- $g \in G$: model belonging to class of interpretable models, e.g. linear model $g(z') = w_g.z'$
- $\Omega(g)$: Complexity of g model
  - Depth of trees in decision trees
  - Number of weights in linear models

LIME: Fidelity-Interpretability Trade-off

Notations:

- $\mathcal{L}(f, g, \pi_x)$: measure of how unfaithful g is in approximating f in the locality defined by π_x
- $g \in G$: model belonging to class of interpretable models, e.g. linear model $g(z') = w_g.z'$
- $\Omega(g)$: Complexity of g model
  - Depth of trees in decision trees
  - Number of weights in linear models
  - For images, $\Omega(g) = \mathbb{1}[||w_g||_0 > K]$ where K is limit on number of super-pixels

LIME: Fidelity-Interpretability Trade-off

Notations:

- $\mathcal{L}(f, g, \pi_x)$: measure of how unfaithful g is in approximating f in the locality defined by π_x
- $g \in G$: model belonging to class of interpretable models, e.g. linear model $g(z') = w_g.z'$
- $\Omega(g)$: Complexity of g model
  - Depth of trees in decision trees
  - Number of weights in linear models
  - For images, $\Omega(g) = \mathbb{1}[||w_g||_0 > K]$ where K is limit on number of super-pixels
- **LIME explanation** obtained as a trade-off:

$$\varepsilon(x) = \arg\min_g \mathcal{L}(f, g, \pi_x) + \Omega(g)$$

SHAP[8]

- Inspired from Shapley values in game theory
- let N: Total number of features; v: Value function that assigns a real number to any coalition $S \subseteq N$; and $\phi_v(i)$: Attribution score for feature i

[8] Lundberg et al, A Unified Approach to Interpreting Model Predictions; NeurIPS 2017

SHAP[8]

- Inspired from Shapley values in game theory
- let N: Total number of features; v: Value function that assigns a real number to any coalition $S \subseteq N$; and $\phi_v(i)$: Attribution score for feature i
- **Attribution score:** Marginal contribution that player (in our case, feature) i makes upon joining the team, averaged over all orders in which team can be formed

$$\phi_v(i) = \sum_{S \subseteq \{1,2,\ldots,N\} \setminus \{i\}} \frac{1}{N!} |S|!(N - |S| - 1)! \underbrace{(v(S \cup i) - v(S))}_{\text{Value of adding player } i \text{ to a coalition}}$$

[8] Lundberg et al, A Unified Approach to Interpreting Model Predictions; NeurIPS 2017

SHAP[8]

- Inspired from Shapley values in game theory
- let N: Total number of features; v: Value function that assigns a real number to any coalition $S \subseteq N$; and $\phi_v(i)$: Attribution score for feature i
- **Attribution score:** Marginal contribution that player (in our case, feature) i makes upon joining the team, averaged over all orders in which team can be formed

$$\phi_v(i) = \sum_{S \subseteq \{1,2,...,N\} \setminus \{i\}} \frac{1}{N!} |S|!(N - |S| - 1)! \underbrace{(v(S \cup i) - v(S))}_{\text{Value of adding player } i \text{ to a coalition}}$$

- With $f(x)$ as model prediction, we marginalize over out-of-coalition features $x_{\bar{S}}$ where $\bar{S} = \{1, 2, ..., N\} \setminus S$ to get:

$$v(S) = \mathbb{E}_{p(x'|x_S)}[f(x_S \cup x'_{\bar{S}})]$$

- SHAP assumes features to be independent $\implies v(S) = \mathbb{E}_{p(x')}[f(x_S \cup x'_{\bar{S}})]$

[8] Lundberg et al, A Unified Approach to Interpreting Model Predictions; NeurIPS 2017

DeepSHAP

- Assumes input features are independent of one another and explanation model is linear

DeepSHAP

- Assumes input features are independent of one another and explanation model is linear
- Take distribution of baselines and compute DeepLIFT attribution for each input-baseline pair, then average resulting attributions per input example

DeepSHAP

- Assumes input features are independent of one another and explanation model is linear
- Take distribution of baselines and compute DeepLIFT attribution for each input-baseline pair, then average resulting attributions per input example

How to Evaluate Explanations?

[9] Melis et al, Towards Robust Interpretability with Self-Explaining Neural Networks, NeurIPS 2018

[10] Petsiuk et al, RISE: Randomized Input Sampling for Explanation of Black-box Models, BMVC 2018

How to Evaluate Explanations?

- IoU of thresholded salient region with ground truth bounding box (if available)

[9] Melis et al, Towards Robust Interpretability with Self-Explaining Neural Networks, NeurIPS 2018
[10] Petsiuk et al, RISE: Randomized Input Sampling for Explanation of Black-box Models, BMVC 2018

How to Evaluate Explanations?

- IoU of thresholded salient region with ground truth bounding box (if available)
- **Faithfulness**[9]: Correlation between attribution scores and output differences on perturbation:

$$F = \langle \rho(R, \Delta) \rangle_{p(\mathbf{x})}$$

where R_i is relevance of pixel i and $\Delta_i = f(\mathbf{x}) - f(\mathbf{x}_i)$ where $\mathbf{x}_i$ is image obtained after perturbing pixel i

[9] Melis et al, Towards Robust Interpretability with Self-Explaining Neural Networks, NeurIPS 2018

[10] Petsiuk et al, RISE: Randomized Input Sampling for Explanation of Black-box Models, BMVC 2018

How to Evaluate Explanations?

- IoU of thresholded salient region with ground truth bounding box (if available)
- **Faithfulness**[9]: Correlation between attribution scores and output differences on perturbation:

$$F = \langle \rho(R, \Delta) \rangle_{p(\mathbf{x})}$$

where R_i is relevance of pixel i and $\Delta_i = f(\mathbf{x}) - f(\mathbf{x}_i)$ where $\mathbf{x}_i$ is image obtained after perturbing pixel i

- **Causal Metric (Deletion Metric)**[10]:
  1. Delete pixels sequentially, most relevant first

[9] Melis et al, Towards Robust Interpretability with Self-Explaining Neural Networks, NeurIPS 2018

[10] Petsiuk et al, RISE: Randomized Input Sampling for Explanation of Black-box Models, BMVC 2018

How to Evaluate Explanations?

- IoU of thresholded salient region with ground truth bounding box (if available)
- **Faithfulness**[9]: Correlation between attribution scores and output differences on perturbation:

$$F = \langle \rho(R, \Delta) \rangle_{p(\mathbf{x})}$$

where R_i is relevance of pixel i and $\Delta_i = f(\mathbf{x}) - f(\mathbf{x}_i)$ where $\mathbf{x}_i$ is image obtained after perturbing pixel i

- **Causal Metric (Deletion Metric)**[10]:
  1. Delete pixels sequentially, most relevant first
  2. Compute AUC of network's output as function of perturbed inputs vs amount of perturbation; lesser AUC better

[9] Melis et al, Towards Robust Interpretability with Self-Explaining Neural Networks, NeurIPS 2018

[10] Petsiuk et al, RISE: Randomized Input Sampling for Explanation of Black-box Models, BMVC 2018

How to Evaluate Explanations?

- IoU of thresholded salient region with ground truth bounding box (if available)
- **Faithfulness**[9]: Correlation between attribution scores and output differences on perturbation:

$$F = \langle \rho(R, \Delta) \rangle_{p(\mathbf{x})}$$

where R_i is relevance of pixel i and $\Delta_i = f(\mathbf{x}) - f(\mathbf{x}_i)$ where $\mathbf{x}_i$ is image obtained after perturbing pixel i

- **Causal Metric (Deletion Metric)**[10]:
  1. Delete pixels sequentially, most relevant first
  2. Compute AUC of network's output as function of perturbed inputs vs amount of perturbation; lesser AUC better

  Similarly, **Insertion Metric** inserts pixels sequentially, least relevant first; higher AUC better

[9] Melis et al, Towards Robust Interpretability with Self-Explaining Neural Networks, NeurIPS 2018

[10] Petsiuk et al, RISE: Randomized Input Sampling for Explanation of Black-box Models, BMVC 2018

How to Evaluate Explanations?

- **ROAR: RemOve And Retrain**[11]:
  1. Get saliency map for each image in training data

[11] Hooker et al, A Benchmark for Interpretability Methods in Deep Neural Networks, NeurIPS 2019
[12] Adebayo et al, Sanity Checks for Saliency Maps, NeurIPS 2018
[13] Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017

How to Evaluate Explanations?

- **ROAR: RemOve And Retrain**[11]:
  1. Get saliency map for each image in training data
  2. Retrain the model after perturbing most relevant pixels

[11] Hooker et al, A Benchmark for Interpretability Methods in Deep Neural Networks, NeurIPS 2019
[12] Adebayo et al, Sanity Checks for Saliency Maps, NeurIPS 2018
[13] Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017

How to Evaluate Explanations?

- **ROAR: RemOve And Retrain**[11]:
  1. Get saliency map for each image in training data
  2. Retrain the model after perturbing most relevant pixels
  3. New model should have large reduction in accuracy

[11] Hooker et al, A Benchmark for Interpretability Methods in Deep Neural Networks, NeurIPS 2019
[12] Adebayo et al, Sanity Checks for Saliency Maps, NeurIPS 2018
[13] Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017

How to Evaluate Explanations?

- **ROAR: RemOve And Retrain**[11]:
  1. Get saliency map for each image in training data
  2. Retrain the model after perturbing most relevant pixels
  3. New model should have large reduction in accuracy
- **Sanity checks for saliency maps**[12](Homework reading!)

[11] Hooker et al, A Benchmark for Interpretability Methods in Deep Neural Networks, NeurIPS 2019
[12] Adebayo et al, Sanity Checks for Saliency Maps, NeurIPS 2018
[13] Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017

How to Evaluate Explanations?

- **ROAR: RemOve And Retrain**[11]:
  1. Get saliency map for each image in training data
  2. Retrain the model after perturbing most relevant pixels
  3. New model should have large reduction in accuracy
- **Sanity checks for saliency maps**[12](Homework reading!)
- **Axioms for attribution**[13](Homework reading!)

[11] Hooker et al, A Benchmark for Interpretability Methods in Deep Neural Networks, NeurIPS 2019
[12] Adebayo et al, Sanity Checks for Saliency Maps, NeurIPS 2018
[13] Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017

Summary

- Both DeepLIFT and Integrated Gradients overcome saturating gradients problem; although DeepLIFT is usually faster, it violates Implementation Invariance axiom[14] (one of the axioms for homework reading!) due to use of discrete gradients
- Smooth Integrated Gradients may be preferred over Integrated Gradients when sparsity is desired
- For better interpretability in terms of visual coherence, XRAI is good choice whose mask is composed of relevant segments rather than pixels
- LIME is model-agnostic and can be used for image, text as well as tabular data but is slow and appears inconsistent between runs
- SHAP has strong game-theoretic background but needs approximations for real world experiments

[14] Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017

Homework

Reading

- Go through list of axioms of attribution in Sundararajan et al, Axiomatic Attribution for Deep Networks, ICML 2017 and for each axiom try to identify the attribution algorithms that satisfy that
- Go through proposed sanity checks and experimental findings in Adebayo et al, Sanity Checks for Saliency Maps, NeurIPS 2018

Programming

- Play with Captum: A popular library for model interpretation by Facebook Open Source
- Try visualizing your models through the lens of OpenAI Microscope

Extra Resources

- Molnar, Interpretable machine learning: A Guide for Making Black Box Models Explainable, 2019: https://christophm.github.io/interpretable-ml-book/.
- For a collection of tutorials and software packages, please refer: https://github.com/jphall663/awesome-machine-learning-interpretability

References I

Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman. "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps". In: *2nd International Conference on Learning Representations, ICLR 2014, Banff, AB, Canada, April 14-16, 2014, Workshop Track Proceedings*. Ed. by Yoshua Bengio and Yann LeCun. 2014.

Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. ""Why Should I Trust You?": Explaining the Predictions of Any Classifier". In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Francisco, CA, USA, August 13-17, 2016*. 2016, pp. 1135–1144.

Scott M Lundberg and Su-In Lee. "A Unified Approach to Interpreting Model Predictions". In: *Advances in Neural Information Processing Systems 30*. Ed. by I. Guyon et al. Curran Associates, Inc., 2017, pp. 4765–4774.

References II

Avanti Shrikumar, Peyton Greenside, and Anshul Kundaje. "Learning Important Features Through Propagating Activation Differences". In: *Proceedings of the 34th International Conference on Machine Learning*. Ed. by Doina Precup and Yee Whye Teh. Vol. 70. Proceedings of Machine Learning Research. International Convention Centre, Sydney, Australia: PMLR, 2017, pp. 3145–3153.

D. Smilkov et al. "SmoothGrad: removing noise by adding noise". In: *ICML workshop on visualization for deep learning* (June 2017). arXiv: 1706.03825 [cs.LG].

Mukund Sundararajan, Ankur Taly, and Qiqi Yan. "Axiomatic Attribution for Deep Networks". In: *Proceedings of the 34th International Conference on Machine Learning*. Ed. by Doina Precup and Yee Whye Teh. Vol. 70. Proceedings of Machine Learning Research. International Convention Centre, Sydney, Australia: PMLR, 2017, pp. 3319–3328.

Julius Adebayo et al. "Sanity Checks for Saliency Maps". In: *Advances in Neural Information Processing Systems 31*. Ed. by S. Bengio et al. Curran Associates, Inc., 2018, pp. 9505–9515.

References III

David Alvarez Melis and Tommi Jaakkola. "Towards Robust Interpretability with Self-Explaining Neural Networks". In: *Advances in Neural Information Processing Systems 31*. Ed. by S. Bengio et al. Curran Associates, Inc., 2018, pp. 7775–7784.

Vitali Petsiuk, Abir Das, and Kate Saenko. "RISE: Randomized Input Sampling for Explanation of Black-box Models". In: *Proceedings of the British Machine Vision Conference (BMVC)*. 2018.

Sara Hooker et al. "A Benchmark for Interpretability Methods in Deep Neural Networks". In: *Advances in Neural Information Processing Systems 32*. Ed. by H. Wallach et al. Curran Associates, Inc., 2019, pp. 9737–9748.

Andrei Kapishnikov et al. "XRAI: Better Attributions Through Regions". In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2019.

http://www.unofficialgoogledatascience.com/2017/03/attributing-deep-networks-prediction-to.html.