

LSTMs and GRUs

Vineeth N Balasubramanian

Department of Computer Science and Engineering
Indian Institute of Technology, Hyderabad



Review: Questions

Question

- How to tackle the vanishing gradient problem in RNNs with solutions that don't change the architecture?


Review: Questions

Question

- How to tackle the vanishing gradient problem in RNNs with solutions that don't change the architecture? **Use ReLU; Regularization; Better initialization of weights; Use only short time sequences**


Long-Term Dependencies: The Problem

- RNNs connect previous information to present task which:
 - may be enough for predicting the next word for “the clouds are in the *sky*”




Long-Term Dependencies: The Problem

- RNNs connect previous information to present task which:
 - may be enough for predicting the next word for “the clouds are in the sky”



- may not be enough when more context is needed: “I grew up in France ... I speak fluent French”



Recall: Training RNNs

- Recommended method: Backprop Through Time (BPTT)

Recall: Training RNNs

- Recommended method: Backprop Through Time (BPTT)
- Limitations of BPTT
 - Vanishing Gradients
 - Exploding Gradients
- How to overcome by changes in RNN architecture?

Recall: Training RNNs

- Recommended method: Backprop Through Time (BPTT)
- Limitations of BPTT
 - Vanishing Gradients
 - Exploding Gradients
- How to overcome by changes in RNN architecture?
 - **LSTMs** (1997)
 - **GRUs** (2014)

Long Short-Term Memory (LSTM)¹

- A type of RNN proposed by Hochreiter and Schmidhuber in 1997 as a solution to the long-term dependency problem

Credit: Christopher Manning, Stanford Univ

¹Hochreiter and Schmidhuber, Long Short-Term Memory, Neural Computation, 1997

Long Short-Term Memory (LSTM)¹

- A type of RNN proposed by Hochreiter and Schmidhuber in 1997 as a solution to the long-term dependency problem
- On step t , there is a **hidden state** $h^{(t)}$ and a **cell state** $c^{(t)}$

Credit: Christopher Manning, Stanford Univ

¹Hochreiter and Schmidhuber, Long Short-Term Memory, Neural Computation, 1997

Long Short-Term Memory (LSTM)¹

- A type of RNN proposed by Hochreiter and Schmidhuber in 1997 as a solution to the long-term dependency problem
- On step t , there is a **hidden state** $h^{(t)}$ and a **cell state** $c^{(t)}$
 - The cell **stores long-term information**

Credit: Christopher Manning, Stanford Univ

¹Hochreiter and Schmidhuber, Long Short-Term Memory, Neural Computation, 1997

Long Short-Term Memory (LSTM)¹

- A type of RNN proposed by Hochreiter and Schmidhuber in 1997 as a solution to the long-term dependency problem
- On step t , there is a **hidden state** $h^{(t)}$ and a **cell state** $c^{(t)}$
 - The cell **stores long-term information**
 - The LSTM can **erase**, **write** and **read** information from the cell

Credit: Christopher Manning, Stanford Univ

¹Hochreiter and Schmidhuber, Long Short-Term Memory, Neural Computation, 1997

Long Short-Term Memory (LSTM)¹

- A type of RNN proposed by Hochreiter and Schmidhuber in 1997 as a solution to the long-term dependency problem
- On step t , there is a **hidden state** $h^{(t)}$ and a **cell state** $c^{(t)}$
 - The cell **stores long-term information**
 - The LSTM can **erase**, **write** and **read** information from the cell
- The selection of which information is erased/written/read is controlled by three corresponding **gates**

Credit: Christopher Manning, Stanford Univ

¹Hochreiter and Schmidhuber, Long Short-Term Memory, Neural Computation, 1997

Long Short-Term Memory (LSTM)¹

- A type of RNN proposed by Hochreiter and Schmidhuber in 1997 as a solution to the long-term dependency problem
- On step t , there is a **hidden state** $h^{(t)}$ and a **cell state** $c^{(t)}$
 - The cell **stores long-term information**
 - The LSTM can **erase**, **write** and **read** information from the cell
- The selection of which information is erased/written/read is controlled by three corresponding **gates**
 - On each timestep, each element of the gates can be open (1), closed (0), or somewhere in-between

Credit: Christopher Manning, Stanford Univ

¹Hochreiter and Schmidhuber, Long Short-Term Memory, Neural Computation, 1997

Long Short-Term Memory (LSTM)¹



- A type of RNN proposed by Hochreiter and Schmidhuber in 1997 as a solution to the long-term dependency problem
- On step t , there is a **hidden state** $h^{(t)}$ and a **cell state** $c^{(t)}$
 - The cell **stores long-term information**
 - The LSTM can **erase**, **write** and **read** information from the cell
- The selection of which information is erased/written/read is controlled by three corresponding **gates**
 - On each timestep, each element of the gates can be open (1), closed (0), or somewhere in-between
 - The gates are dynamic, their value is computed based on current context

Credit: Christopher Manning, Stanford Univ

¹Hochreiter and Schmidhuber, Long Short-Term Memory, Neural Computation, 1997

LSTMs



- All RNNs have the form of a **chain of repeating modules** of neural network
- Repeating module in a vanilla RNN is a single layer with \tanh activation



Credit: Christopher Olah


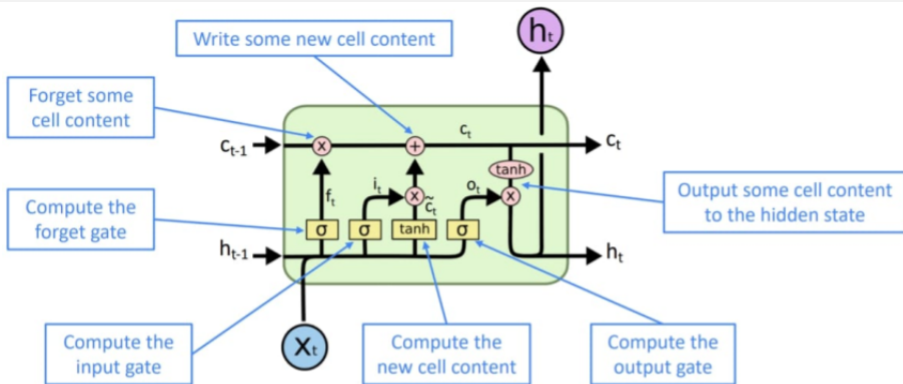
LSTMs

- All RNNs have the form of a **chain of repeating modules** of neural network
- Repeating module in an **LSTM** contains four interacting layers



Credit: Christopher Olah


LSTMs



Credit: Christopher Manning, Stanford Univ

LSTMs Explained


- **Cell state (C_t):**



Credit: Christopher Olah

LSTMs Explained


- **Cell state (C_t):**
 - Information can flow along cell state unchanged. Why is this important?



Credit: Christopher Olah

LSTMs Explained



- **Cell state (C_t):**
 - Information can flow along cell state unchanged. Why is this important?
 - Ability to **remove** or **add** information to cell state, regulated by gates



Credit: Christopher Olah

LSTMs Explained


- **Cell state (C_t):**
 - Information can flow along cell state unchanged. Why is this important?
 - Ability to **remove** or **add** information to cell state, regulated by gates
- **Gates:**



Credit: Christopher Olah

LSTMs Explained


- **Cell state (C_t):**
 - Information can flow along cell state unchanged. Why is this important?
 - Ability to **remove** or **add** information to cell state, regulated by gates
- **Gates:**
 - Composed of a **sigmoid neural net layer** and a **pointwise multiplication operation**



Credit: Christopher Olah

LSTMs Explained


- **Cell state (C_t):**
 - Information can flow along cell state unchanged. Why is this important?
 - Ability to **remove** or **add** information to cell state, regulated by gates
- **Gates:**
 - Composed of a **sigmoid neural net layer** and a **pointwise multiplication operation**
 - **Sigmoid layer** outputs numbers between zero and one, describing how much of each component should be let through



Credit: Christopher Olah

LSTMs Explained

Forget gate: controls what is kept vs what is forgotten, from previous cell state




$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

Credit: Christopher Olah

LSTMs Explained

Input gate: decides what information to throw away from the cell state

Cell content: new content to be written to cell




$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Credit: Christopher Olah

LSTMs Explained

Cell state: erase (“forget”) some content from last cell state, and write (“input”) some new cell content




$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Credit: Christopher Olah

LSTMs Explained

Output gate: controls what parts of cell are output to hidden state

Hidden state: read (“output”) some content from cell




$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$


$$h_t = o_t * \tanh (C_t)$$

Credit: Christopher Olah

LSTMs Explained




LSTMs Explained




- What can you tell about cell state(C_t), if forget gate is set to 1 and input gate set to 0?

LSTMs Explained




- What can you tell about cell state(C_t), if forget gate is set to 1 and input gate set to 0?
 - Information of that cell is **preserved indefinitely**

LSTMs Explained



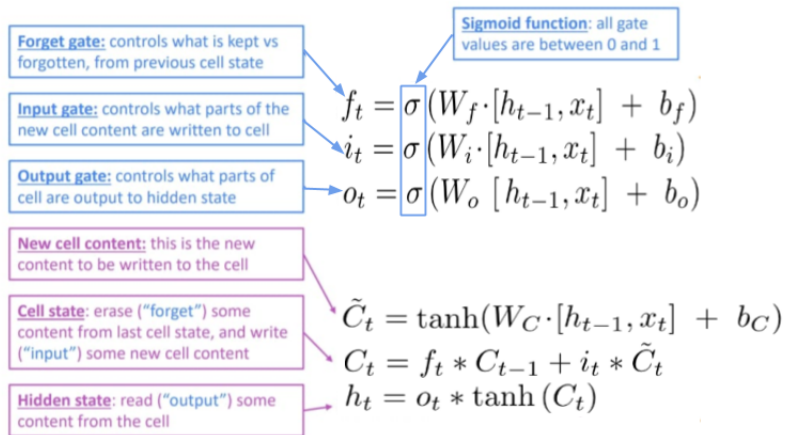
- What can you tell about cell state(C_t), if forget gate is set to 1 and input gate set to 0?
 - Information of that cell is **preserved indefinitely**
- What happens if you fix input gate to all 1s, forget gate to all 0s, output gate to all 1s?

LSTMs Explained




- What can you tell about cell state(C_t), if forget gate is set to 1 and input gate set to 0?
 - Information of that cell is **preserved indefinitely**
- What happens if you fix input gate to all 1s, forget gate to all 0s, output gate to all 1s?
 - Almost **standard RNN**;

LSTMs Explained



- What can you tell about cell state(C_t), if forget gate is set to 1 and input gate set to 0?
 - Information of that cell is **preserved indefinitely**
- What happens if you fix input gate to all 1s, forget gate to all 0s, output gate to all 1s?
 - Almost **standard RNN**; Why almost?


LSTMs Explained




- What can you tell about cell state(C_t), if forget gate is set to 1 and input gate set to 0?
 - Information of that cell is **preserved indefinitely**
- What happens if you fix input gate to all 1s, forget gate to all 0s, output gate to all 1s?
 - Almost **standard RNN**; Why almost?
 - **Tanh** added here

Credit: Christopher Olah; Christopher Manning, Stanford University


LSTM: How does it solve the vanishing gradient problem?




- Gradient “highway”




LSTM: How does it solve the vanishing gradient problem?



- Gradient “highway”
- Gradient at C_t passed on to C_{t-1} unaffected by any other operations, but for forget gate; why does this not matter?




LSTM: How does it solve the vanishing gradient problem?



- Gradient “highway”
- Gradient at C_t passed on to C_{t-1} unaffected by any other operations, but for forget gate; why does this not matter?
- Forget gate is part of the design, it reduces the gradient where it should, does not ameliorate the gradient otherwise!

Variants of LSTM

- **LSTM with peephole connections**²



$$f_t = \sigma(W_f \cdot [C_{t-1}, h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma(W_i \cdot [C_{t-1}, h_{t-1}, x_t] + b_i)$$

$$o_t = \sigma(W_o \cdot [C_t, h_{t-1}, x_t] + b_o)$$


Credit: Christopher Olah

²Gers and Schmidhuber, Recurrent nets that time and count, IJCNN 2000

Variants of LSTM

- **Coupled forget and input gates**

- Instead of separately deciding what to **forget** and what to **add**, make decisions together



$$C_t = f_t * C_{t-1} + (1 - f_t) * \tilde{C}_t$$

Credit: Christopher Olah

History of LSTMs

- **1997 - RTRL + BPTT (No forget gate)**
 - Hochreiter and Schmidhuber, Long Short- Term Memory, Neural Computation, 1997

History of LSTMs

- **1997 - RTRL + BPTT (No forget gate)**
 - Hochreiter and Schmidhuber, Long Short- Term Memory, Neural Computation, 1997
- **1999 – Introduced forget gate**
 - Gers Schmidhuber and Cummins, Learning to forget: Continual prediction with LSTM, ICANN 1999

History of LSTMs

- **1997 - RTRL + BPTT (No forget gate)**
 - Hochreiter and Schmidhuber, Long Short- Term Memory, Neural Computation, 1997
- **1999 – Introduced forget gate**
 - Gers Schmidhuber and Cummins, Learning to forget: Continual prediction with LSTM, ICANN 1999
- **2000 – Peephole connections**
 - Gers and Schmidhuber, Recurrent nets that time and count, IJCNN 2000

History of LSTMs

- **1997 - RTRL + BPTT (No forget gate)**
 - Hochreiter and Schmidhuber, Long Short- Term Memory, Neural Computation, 1997
- **1999 – Introduced forget gate**
 - Gers Schmidhuber and Cummins, Learning to forget: Continual prediction with LSTM, ICANN 1999
- **2000 – Peephole connections**
 - Gers and Schmidhuber, Recurrent nets that time and count, IJCNN 2000
- **2005 – Vanilla LSTM (as we know today) – Used BPTT**
 - Graves and Schmidhuber, Framewise phoneme classification with bidirectional LSTM and other neural network architectures, Neural Networks, 2005

LSTMs: Real-world success

- **2013-2015:** LSTMs started achieving **state-of-the-art results**
 - Successful tasks include: handwriting recognition, speech recognition, machine translation, parsing, image captioning

LSTMs: Real-world success

- **2013-2015**: LSTMs started achieving **state-of-the-art results**
 - Successful tasks include: handwriting recognition, speech recognition, machine translation, parsing, image captioning
- **Now (2020)**, other approaches (e.g. **Transformers**) have become more dominant for certain tasks
 - Transformers use the idea of self-attention
 - In **WMT 2019** ((a MT conference + competition), summary report contains “RNN” 7 times, “Transformer” 105 times

Credit: Christopher Manning, Stanford University

Gated Recurrent Unit (GRU)²

- Proposed in 2014 as a simpler alternative to LSTM

²Chung et al, Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling, NeurIPS-W 2014

Gated Recurrent Unit (GRU)²

- Proposed in 2014 as a simpler alternative to LSTM
- Combines **forget** and **input** gates into a single **update gate**

²Chung et al, Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling, NeurIPS-W 2014


Gated Recurrent Unit (GRU)²

- Proposed in 2014 as a simpler alternative to LSTM
- Combines **forget** and **input** gates into a single **update gate**
- Merges **cell state** and **hidden state**

²Chung et al, Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling, NeurIPS-W 2014

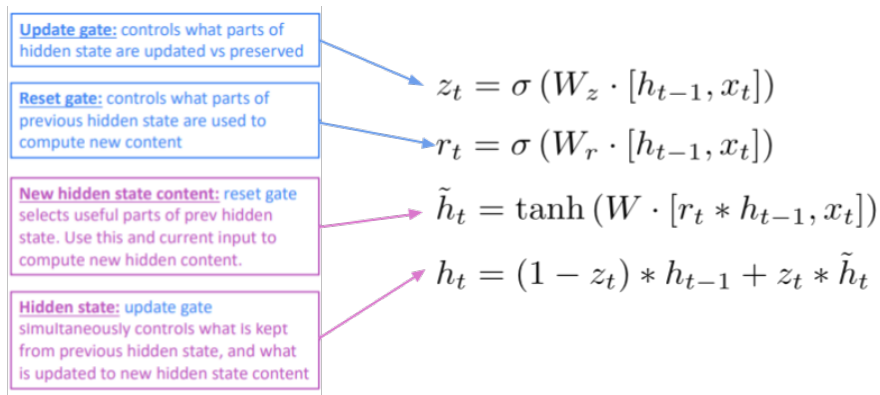
Gated Recurrent Unit (GRU)²

- Proposed in 2014 as a simpler alternative to LSTM
- Combines **forget** and **input** gates into a single **update gate**
- Merges **cell state** and **hidden state**




²Chung et al, Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling, NeurIPS-W 2014

Gated Recurrent Unit (GRU)²



²Chung et al, Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling, NeurIPS-W 2014

Gated Recurrent Unit (GRU)²



- What happens if reset gate is set to all 1s and update gate to all 0s?

²Chung et al, Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling, NeurIPS-W 2014

GRUs vs LSTMs: Summary

- Input and forget gates of LSTMs are coupled by an update gate in GRUs; reset gate (GRUs) is applied directly to previous hidden state

GRUs vs LSTMs: Summary

- Input and forget gates of LSTMs are coupled by an update gate in GRUs; reset gate (GRUs) is applied directly to previous hidden state
- GRU has **two gates**, an LSTM has **three gates**; what does this tell you?

GRUs vs LSTMs: Summary

- Input and forget gates of LSTMs are coupled by an update gate in GRUs; reset gate (GRUs) is applied directly to previous hidden state
- GRU has **two gates**, an LSTM has **three gates**; what does this tell you? **Lesser parameters to learn!**

GRUs vs LSTMs: Summary

- Input and forget gates of LSTMs are coupled by an update gate in GRUs; reset gate (GRUs) is applied directly to previous hidden state
- GRU has **two gates**, an LSTM has **three gates**; what does this tell you? **Lesser parameters to learn!**
- In GRUs:
 - No internal memory (c_t) different from exposed hidden state
 - No output gate as in LSTMs

GRUs vs LSTMs: Summary

- Input and forget gates of LSTMs are coupled by an update gate in GRUs; reset gate (GRUs) is applied directly to previous hidden state
- GRU has **two gates**, an LSTM has **three gates**; what does this tell you? **Lesser parameters to learn!**
- In GRUs:
 - No internal memory (c_t) different from exposed hidden state
 - No output gate as in LSTMs
- LSTM a good default choice (especially if data has long-range dependencies, or if training data is large); Switch to GRUs for speed and fewer parameters

Homework






Readings

- Deep Learning book: Sections 10.1-10.7, 10.10-10.11
- Understanding LSTM Networks
- Illustrated Guide to LSTMs and GRUs: A step by step explanation
- (Optional) Recurrent Neural Network Tutorial– Implementing a GRU/LSTM RNN with Python and Theano
- (Optional) Training LSTMs using BPTT: [Alex Graves' book on RNN](#) (Sec 4.6, pg 36-38)

Questions

- How does GRU address vanishing gradients?

References

-  S. Hochreiter and J. Schmidhuber. “Long Short-Term Memory”. In: *Neural Computation* 9 (1997), pp. 1735–1780.
-  Razvan Pascanu, Tomas Mikolov, and Yoshua Bengio. “On the difficulty of training recurrent neural networks”. In: *ICML*. 2013.
-  Junyoung Chung et al. “Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling”. In: *CoRR* abs/1412.3555 (2014). arXiv: [1412.3555](https://arxiv.org/abs/1412.3555).
-  Li, Fei-Fei; Johnson, Justin; Serena, Yeung; CS 231n - *Convolutional Neural Networks for Visual Recognition* (Spring 2019). URL: <http://cs231n.stanford.edu/2019/> (visited on 08/28/2020).
-  Manning, Christopher, CS 224n *Natural Language Processing with Deep Learning* (Winter 2019). URL: <http://cs224n.stanford.edu/> (visited on 08/28/2020).