

Deep Generative Models: Image Applications

Vineeth N Balasubramanian

Department of Computer Science and Engineering
Indian Institute of Technology, Hyderabad



Homework

Question: Why is MI Gap and not MI used as a metric for disentanglement?

Homework

Question: Why is MI Gap and not MI used as a metric for disentanglement?

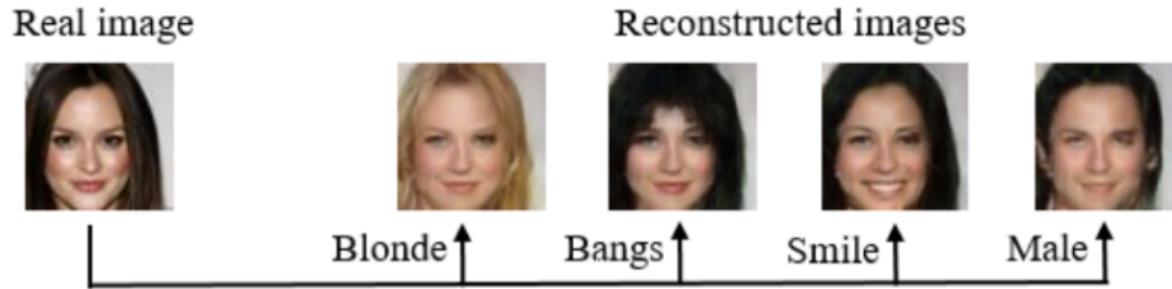
- MI Gap penalizes unaligned latent variables (contains information about more than one generative factor) → undesirable for disentanglement
- If one latent variable reliably models a generative factor, not required for other latent variables to be informative about this factor
- Read Chen et al, Isolating Sources of Disentanglement in Variational Autoencoders, NeurIPS 2018 for more information!

GANs for Image Editing

- Simple image edits such as grayscaling, brightness adjustment or contrast don't require an understanding of the image

GANs for Image Editing

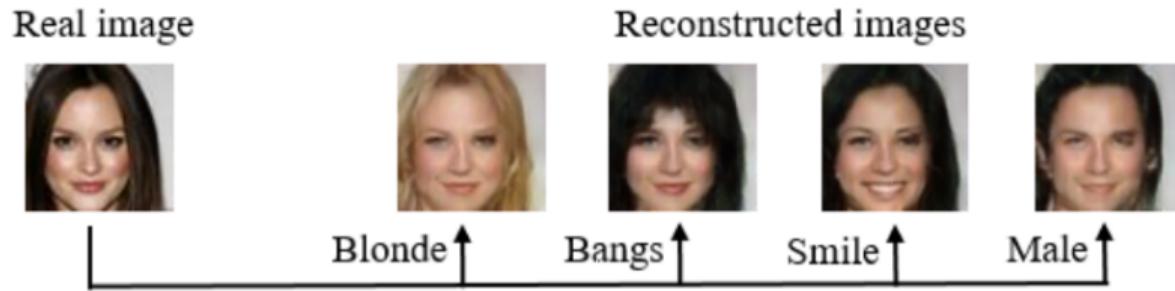
- Simple image edits such as grayscaling, brightness adjustment or contrast don't require an understanding of the image



- However, challenging operations such as changing attributes of a face require sound understanding of the input image. Can GANs help?

GANs for Image Editing

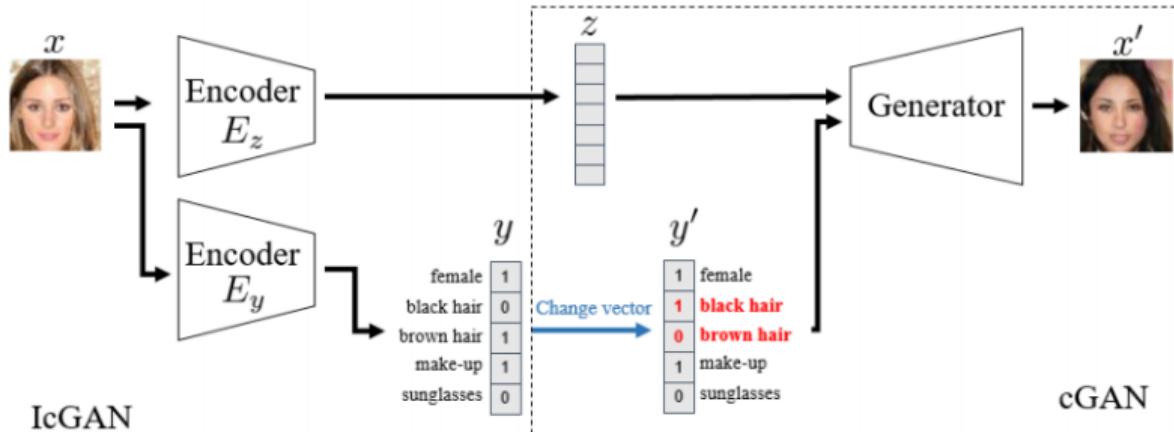
- Simple image edits such as grayscaling, brightness adjustment or contrast don't require an understanding of the image



- However, challenging operations such as changing attributes of a face require sound understanding of the input image. Can GANs help?
- But a GAN on its own does not have the mechanism to map a real image to its latent representation

Image Editing with Invertible Conditional GANs (IcGAN)¹

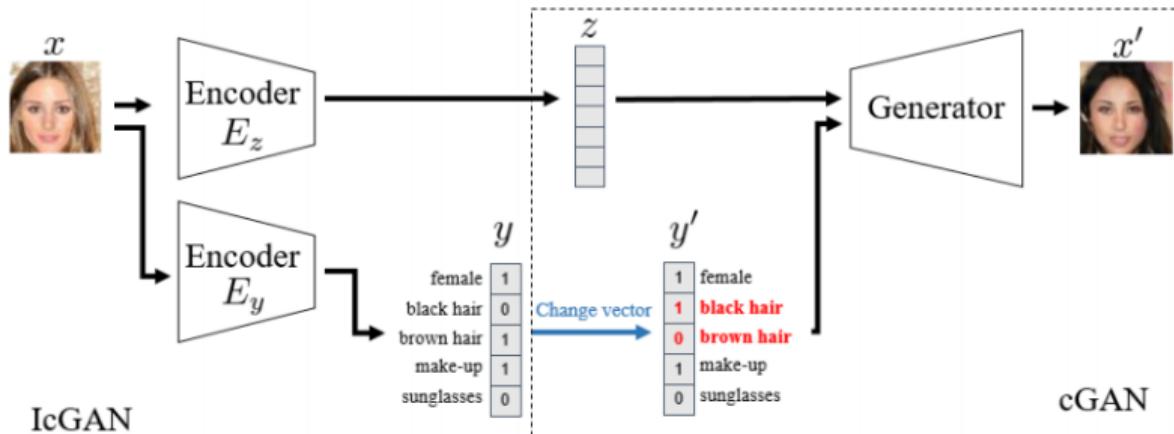
- Combines a conditional GAN with an encoder which can encode an image to its latent representation



¹Perarnau et al, Invertible Conditional GANs for Image Editing, NeurIPS-W 2016

Image Editing with Invertible Conditional GANs (IcGAN)¹

- Combines a conditional GAN with an encoder which can encode an image to its latent representation

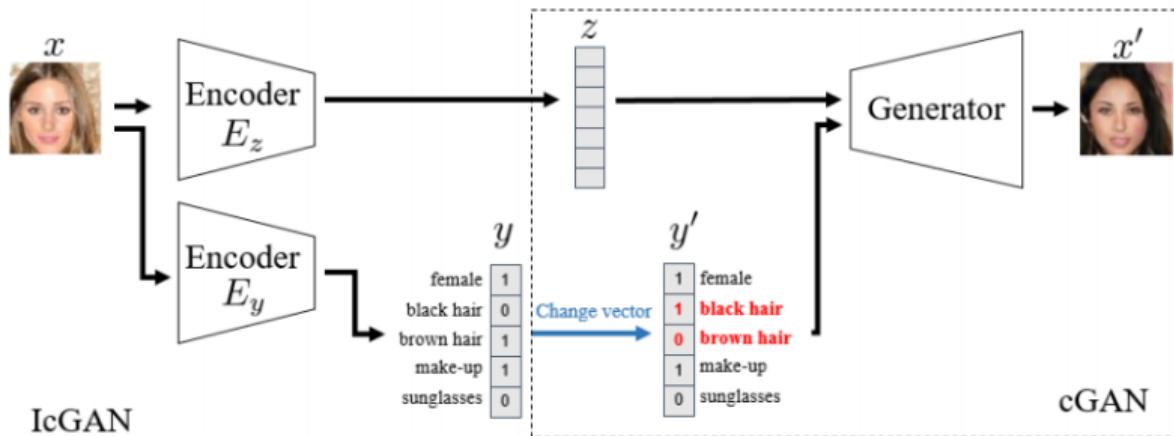


- Can be used for face editing as follows: (1) Pass image through encoder; (2) Change attribute vector y ; (3) Pass through generator

¹Perarnau et al, Invertible Conditional GANs for Image Editing, NeurIPS-W 2016

Invertible Conditional GANs

- Generator G samples an image x' from a latent representation z and conditional information y i.e., $G(z, y) = x'$



- Encoder performs the inverse i.e., $E(x) = (z, y)$

Training IcGAN

- First, a conditional GAN is trained to optimize the following objective:

$$\min_g \max_d v(\theta_g, \theta_d) = E_{x,y \sim p_{data}} [\log D(x, y)] + E_{z \sim p_z. x \sim p_x} [\log(1 - D(G(z, y'), y'))]$$

Training IcGAN

- First, a conditional GAN is trained to optimize the following objective:

$$\min_g \max_d v(\theta_g, \theta_d) = E_{x,y \sim p_{data}} [\log D(x, y)] + E_{z \sim p_z, x \sim p_x} [\log(1 - D(G(z, y'), y'))]$$

- Encoder has two parts: E_z which maps an image to its latent representation, and E_y which maps an image to its conditional information (attributes)

Training IcGAN

- First, a conditional GAN is trained to optimize the following objective:

$$\min_g \max_d v(\theta_g, \theta_d) = E_{x,y \sim p_{data}} [\log D(x, y)] + E_{z \sim p_z, x \sim p_x} [\log(1 - D(G(z, y'), y'))]$$

- Encoder has two parts: E_z which maps an image to its latent representation, and E_y which maps an image to its conditional information (attributes)
- To train E_z , generator is used to generate a dataset of images. Pairs of images and their latent representations (x', z) are used to train E_z using reconstruction loss:

$$\mathcal{L}_{ez} = E_{z \sim p_z, y' \sim p_y} \|z - E_z(G(z, y'))\|_2^2$$

Training IcGAN

- First, a conditional GAN is trained to optimize the following objective:

$$\min_g \max_d v(\theta_g, \theta_d) = E_{x,y \sim p_{data}} [\log D(x, y)] + E_{z \sim p_z, x \sim p_x} [\log(1 - D(G(z, y'), y'))]$$

- Encoder has two parts: E_z which maps an image to its latent representation, and E_y which maps an image to its conditional information (attributes)
- To train E_z , generator is used to generate a dataset of images. Pairs of images and their latent representations (x', z) are used to train E_z using reconstruction loss:

$$\mathcal{L}_{ez} = E_{z \sim p_z, y' \sim p_y} \|z - E_z(G(z, y'))\|_2^2$$

- E_y is trained using real image and attribute pairs:

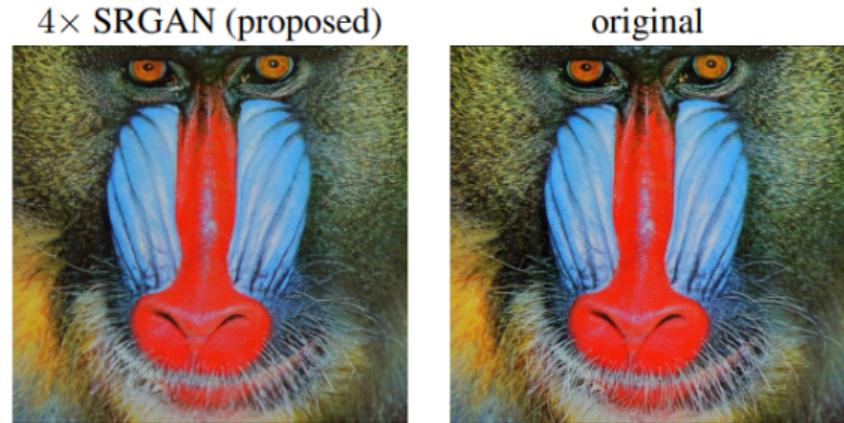
$$\mathcal{L}_{ey} = E_{x,y \sim p_{data}} \|y - E_y(x)\|_2^2$$

IcGAN: Results on CelebA dataset



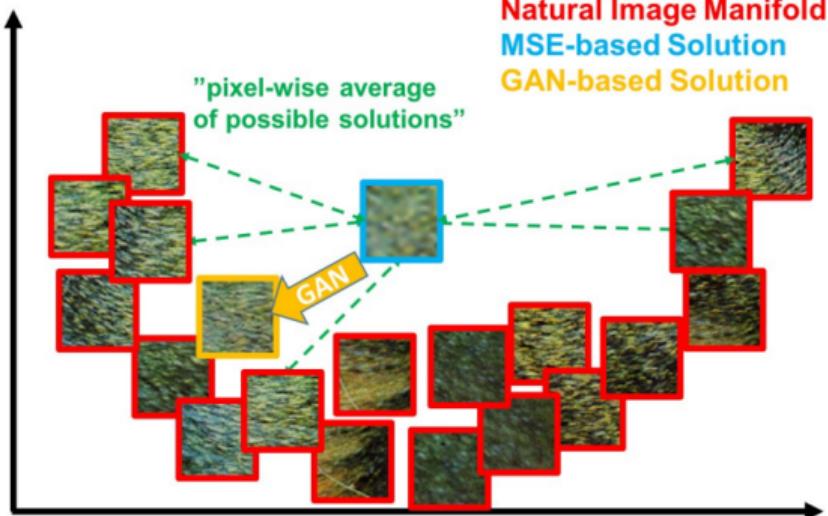
GANs for Super-Resolution²

- **Image Super-Resolution:** Task of obtaining a High-Resolution (HR) image from a Low Resolution (LR) image, a challenging task
- **SRGAN (Super-Resolution GAN)** aims to generate photo-realistic images with $4\times$ upscaling factors



²Ledig et al, Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network, CVPR 2017

SR-GAN



- **Super-Resolution CNNs** minimize Mean Square Error (MSE) → tend to overly smoothen output images, due to pixel-wise averaging of outcomes
- **SR-GAN**, because of its adversarial objective, drives outputs to natural image manifold → results in high quality super-resolved images

SR-GAN: Content and Adversarial Losses

- Uses a pretrained VGG-19 network to compute **content loss** as MSE between VGG-19 feature representations of generator ($G_{\theta_G}(I^{LR})$) and true image I^{HR} :

$$l_{VGG/i.j}^{SR} = \frac{1}{W_{i,j} H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} (\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}(G_{\theta_G}(I^{LR}))_{x,y})^2$$

where $\phi_{i,j}$ = feature map obtained by j-th convolution before i-th max pooling layer, $W_{i,j}$ and $H_{i,j}$ are width and height of feature map respectively

SR-GAN: Content and Adversarial Losses

- Uses a pretrained VGG-19 network to compute **content loss** as MSE between VGG-19 feature representations of generator ($G_{\theta_G}(I^{LR})$) and true image I^{HR} :

$$l_{VGG/i.j}^{SR} = \frac{1}{W_{i,j} H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} (\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}(G_{\theta_G}(I^{LR}))_{x,y})^2$$

where $\phi_{i,j}$ = feature map obtained by j-th convolution before i-th max pooling layer, $W_{i,j}$ and $H_{i,j}$ are width and height of feature map respectively

- **Adversarial loss** given by:

$$l_{Gen}^{SR} = \sum_{n=1}^N -\log D_{\theta_D}(G_{\theta_G}(I^{LR}))$$

SR-GAN: Content and Adversarial Losses

- Uses a pretrained VGG-19 network to compute **content loss** as MSE between VGG-19 feature representations of generator ($G_{\theta_G}(I^{LR})$) and true image I^{HR} :

$$l_{VGG/i.j}^{SR} = \frac{1}{W_{i,j}H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} (\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}(G_{\theta_G}(I^{LR}))_{x,y})^2$$

where $\phi_{i,j}$ = feature map obtained by j-th convolution before i-th max pooling layer, $W_{i,j}$ and $H_{i,j}$ are width and height of feature map respectively

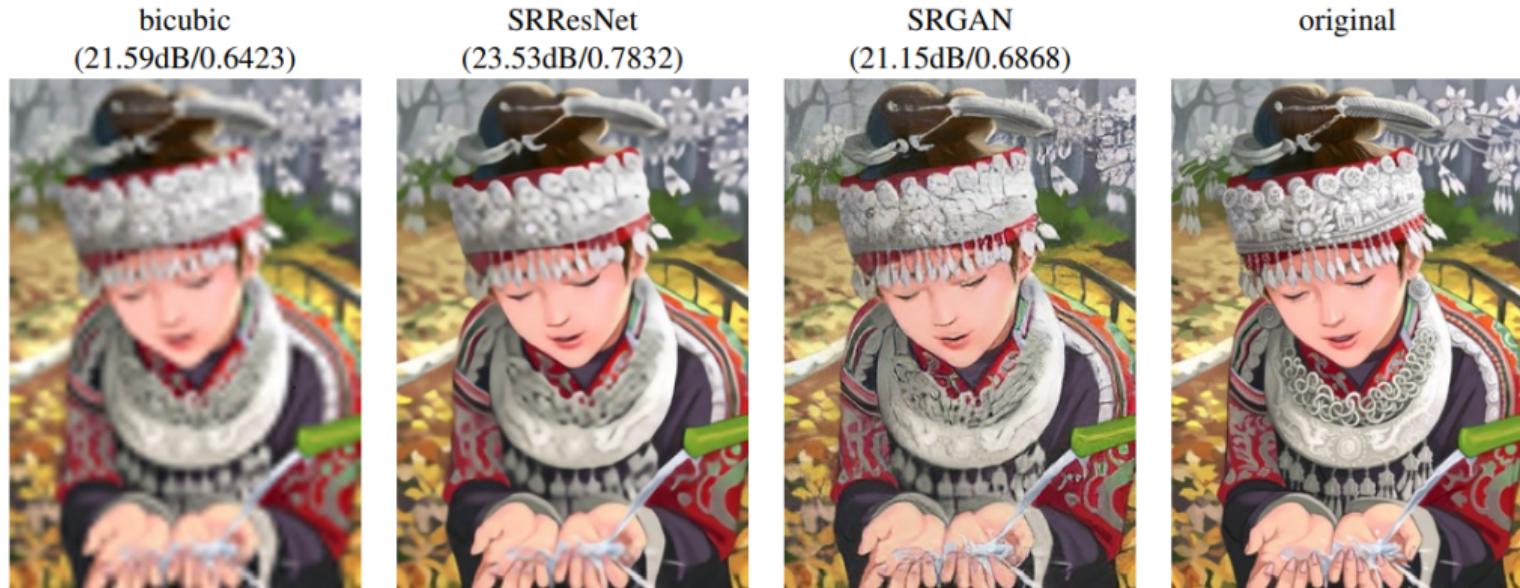
- **Adversarial loss** given by:

$$l_{Gen}^{SR} = \sum_{n=1}^N -\log D_{\theta_D}(G_{\theta_G}(I^{LR}))$$

- **Perceptual loss**, a weighted sum of content and adversarial loss:

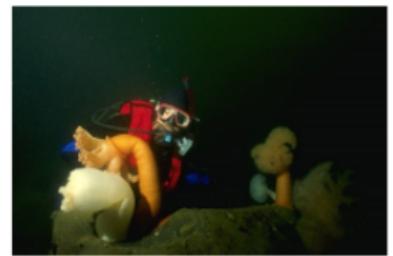
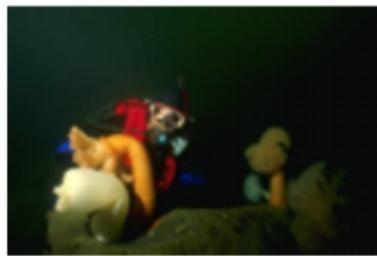
$$l^{SR} = l_X^{SR} + 10^{-3} l_{Gen}^{SR}$$

SR-GAN Results³



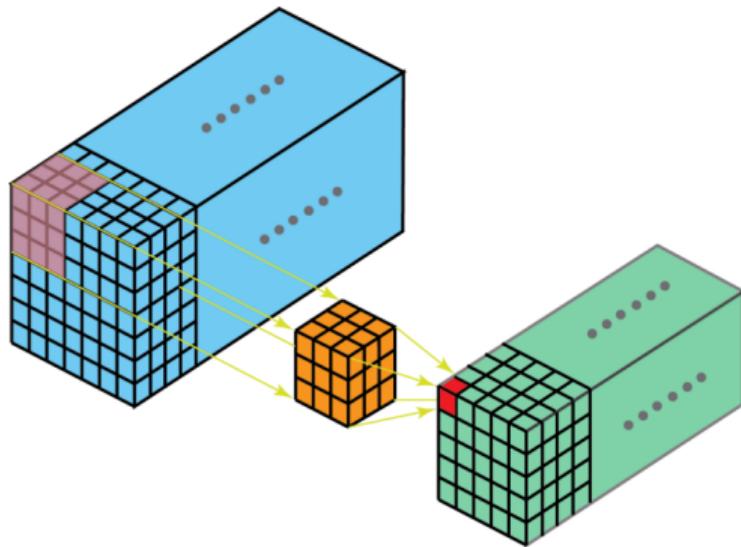
³Ledig et al, Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network, CVPR 2017

SR-GAN More Results⁴



⁴Ledig et al, Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network, CVPR 2017

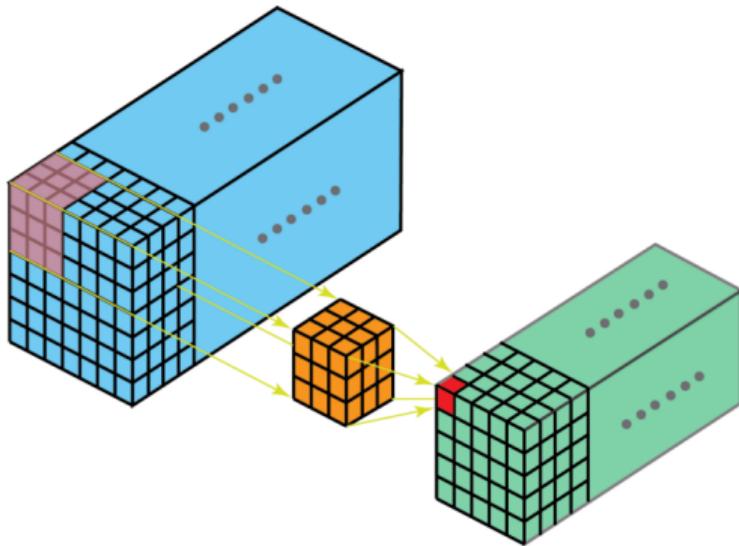
GANs for 3D Object Generation⁵



- We have seen GANs generate photo-realistic 2D images; can GANs generate 3D objects too?

⁵Wu et al, Learning a Probabilistic Latent Space of Object Shapes via 3D Generative Adversarial Modeling, NeurIPS 2016

GANs for 3D Object Generation⁵

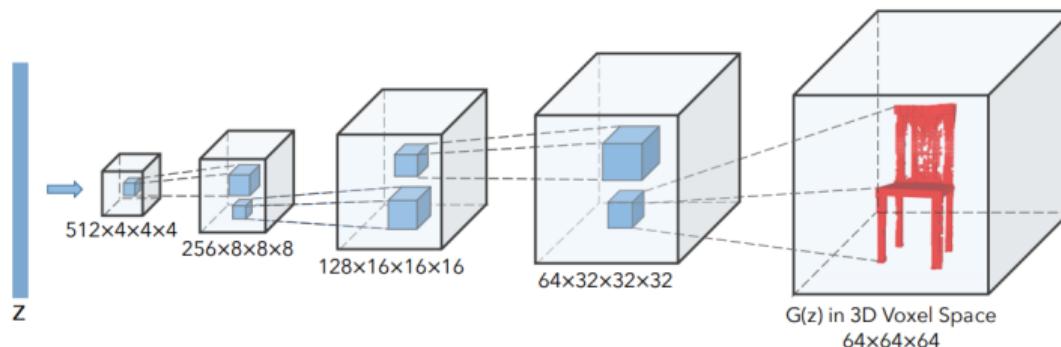


- We have seen GANs generate photo-realistic 2D images; can GANs generate 3D objects too?
- Yes! Recent advances in volumetric (3D) convolutional networks have made it possible to learn 3D object representations

⁵Wu et al, Learning a Probabilistic Latent Space of Object Shapes via 3D Generative Adversarial Modeling, NeurIPS 2016

GANs for 3D Object Generation⁶

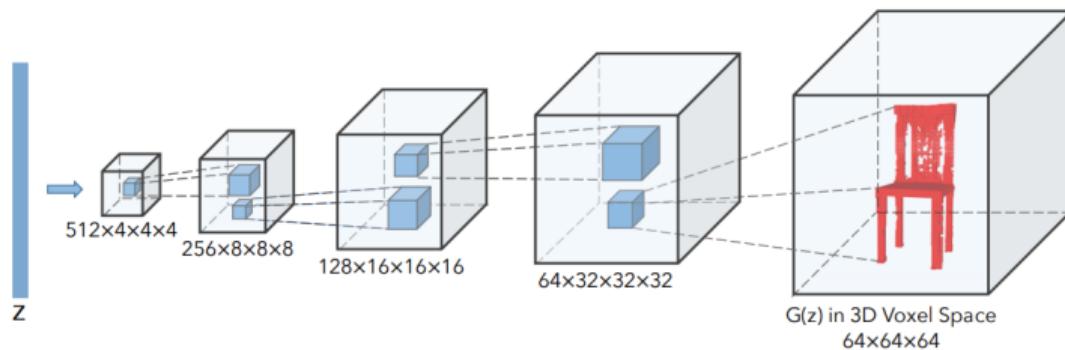
- Generator G maps a 200-length latent vector z to a $64 \times 64 \times 64$ cube, representing an object $G(z)$ in 3D-voxel space



⁶Wu et al, Learning a Probabilistic Latent Space of Object Shapes via 3D Generative Adversarial Modeling, NeurIPS 2016

GANs for 3D Object Generation⁶

- Generator G maps a 200-length latent vector z to a $64 \times 64 \times 64$ cube, representing an object $G(z)$ in 3D-voxel space

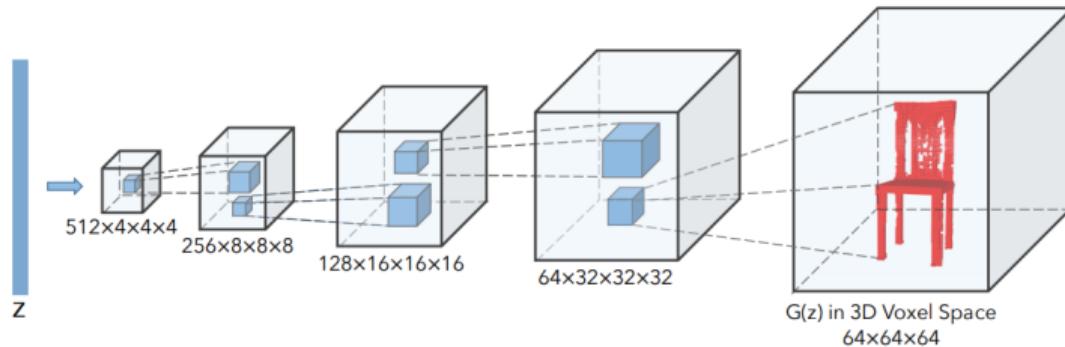


- Discriminator D classifies whether an input 3D object x is real or synthetic

⁶Wu et al, Learning a Probabilistic Latent Space of Object Shapes via 3D Generative Adversarial Modeling, NeurIPS 2016

GANs for 3D Object Generation⁶

- Generator G maps a 200-length latent vector z to a $64 \times 64 \times 64$ cube, representing an object $G(z)$ in 3D-voxel space



- Discriminator D classifies whether an input 3D object x is real or synthetic
- Adversarial loss in 3D-GAN:

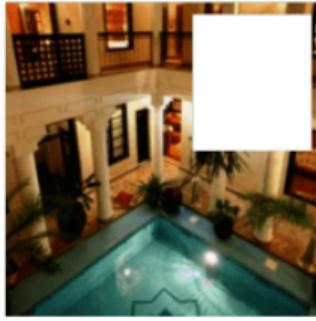
$$L_{3DGAN} = \log D(x) + \log(1 - D(G(z)))$$

⁶Wu et al, Learning a Probabilistic Latent Space of Object Shapes via 3D Generative Adversarial Modeling, NeurIPS 2016

3D-GAN Generated Objects



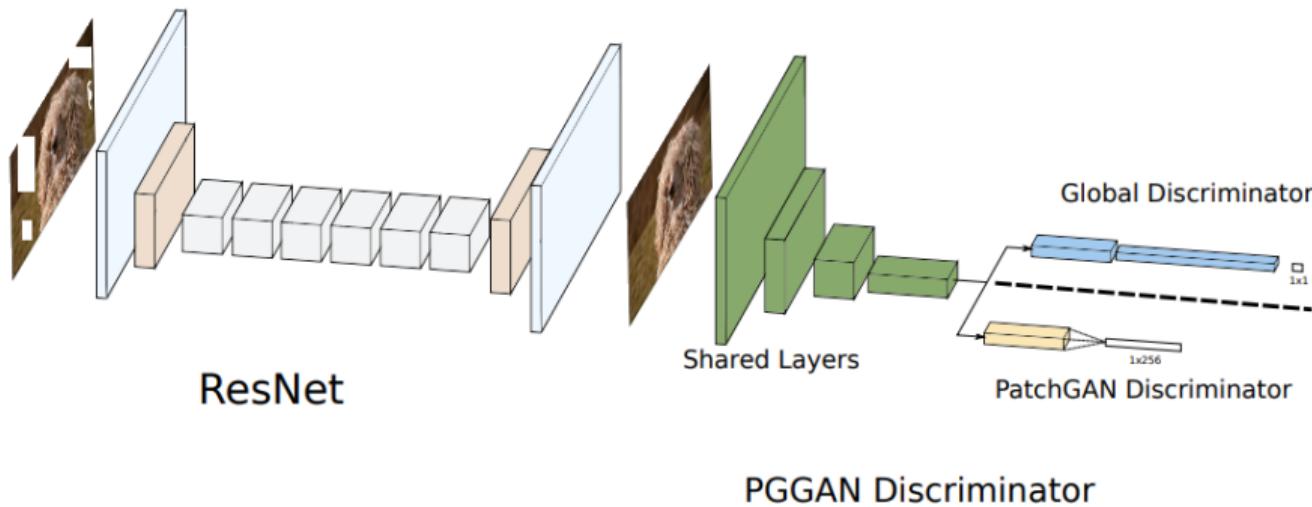
GANs for Image Inpainting⁷



- **Image inpainting** a reconstruction technique used for filling missing parts in an image
- Can a GAN be trained to perform image inpainting?

⁷Demir and Unal, Patch-Based Image Inpainting with Generative Adversarial Networks, arXiv 2018

PG-GAN⁸



- Consists of a ResNet-based generator and a novel discriminator with two heads:
 - **G-GAN discriminator:** reasons globally at image level (decide real vs fake)
 - **Patch GAN discriminator:** reasons locally at patch level (decide real vs fake) - helps capture local texture details

⁸Demir and Unal, Patch-Based Image Inpainting with Generative Adversarial Networks, arXiv 2018

PG-GAN Objective

- **Reconstruction loss:** pixel-wise L1 distance between generated image and ground truth:

$$\mathcal{L}_{rec} = \frac{1}{N} \sum_{i=1}^N \frac{1}{WHC} \|y - x\|_1$$

- **Adversarial loss:** standard GAN objective:

$$\mathcal{L}_{GAN}(G, D) = E_{x \sim p(x)}[\log D(x)] + E_{y \sim p_G(x')}[\log(1 - D(G(x')))]$$

where \mathcal{L}_{g_adv} = adversarial loss for global GAN; \mathcal{L}_{p_adv} = adversarial loss for Patch GAN

- Overall objective:

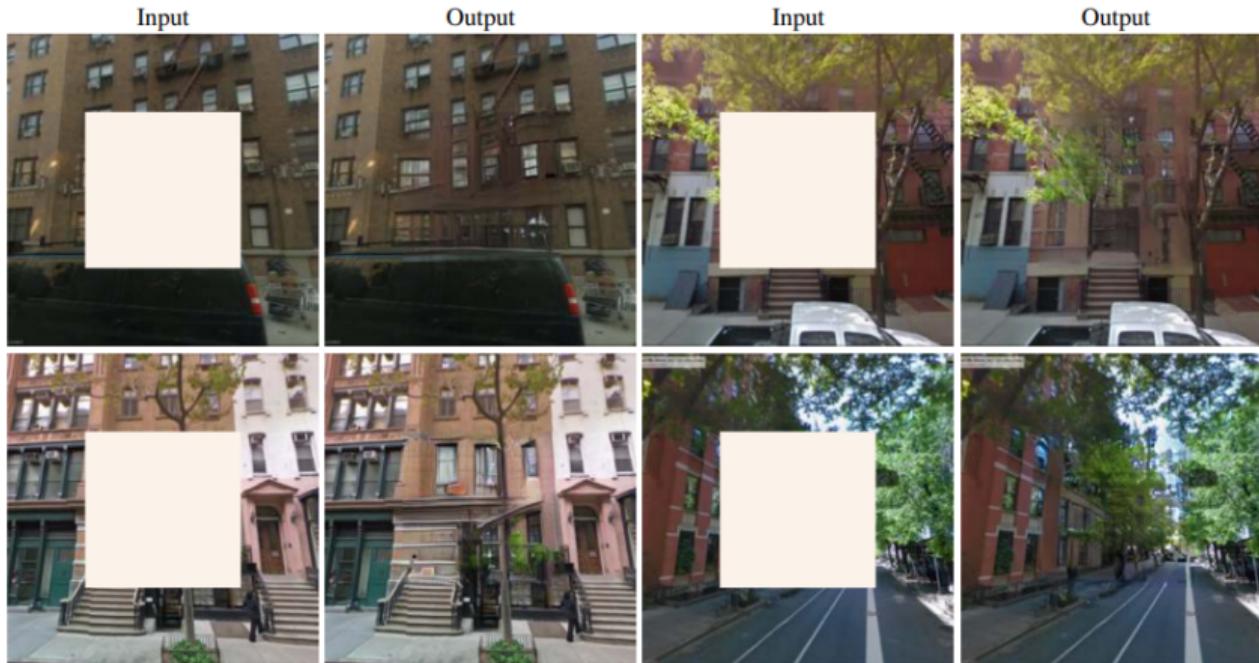
$$\mathcal{L} = \lambda_1 \mathcal{L}_{rec} + \lambda_2 \mathcal{L}_{g_adv} + \lambda_3 \mathcal{L}_{p_adv}$$

PG-GAN Results⁹



⁹Demir and Unal, Patch-Based Image Inpainting with Generative Adversarial Networks, arXiv 2018

PG-GAN More Results¹⁰



¹⁰Demir and Unal, Patch-Based Image Inpainting with Generative Adversarial Networks, arXiv 2018

Homework

Readings

- 3D-GANs
- Avinash H, GAN Zoo
- (Optional) Respective papers

References

-  Guim Perarnau et al. "Invertible Conditional GANs for image editing". In: *ArXiv* abs/1611.06355 (2016).
-  Jiajun Wu et al. "Learning a probabilistic latent space of object shapes via 3d generative-adversarial modeling". In: *Advances in Neural Information Processing Systems*. 2016, pp. 82–90.
-  C. Ledig et al. "Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network". In: *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2017, pp. 105–114.
-  Ugur Demir and Gözde B. Ünal. "Patch-Based Image Inpainting with Generative Adversarial Networks". In: *CoRR* abs/1803.07422 (2018). arXiv: [1803.07422](https://arxiv.org/abs/1803.07422).