



Model Free Prediction : Additional Topics

Easwar Subramanian

TCS Innovation Labs, Hyderabad

Email : easwar.subramanian@tcs.com / cs5500.2020@iith.ac.in

September 27, 2021

Overview



Review

2 Certainity Equivalence Estimate

3 Off Policy Learning



Review



Multi-step TD



▶ One-step TD

$$V^{\pi}(s) = \mathbb{E}_{\pi} \left[r_{t+1} + \gamma V^{\pi}(s_{t+1}) | s_t = s \right]$$

$$V(s_t) \leftarrow V(s_t) + \alpha_t [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

► Two-step TD

$$V^{\pi}(s) = \mathbb{E}_{\pi} \left[r_{t+1} + \gamma r_{t+2} + \gamma^{2} V^{\pi}(s_{t+2}) | s_{t} = s \right]$$

$$V(s_{t}) \leftarrow V(s_{t}) + \alpha_{t} \left[r_{t+1} + \gamma r_{t+2} + \gamma^{2} V(s_{t+2}) - V(s_{t}) \right]$$

 \blacktriangleright More generally, define the n-step return

$$G_t^{(n)} \stackrel{\text{def}}{=} r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n V(s_{t+n})$$

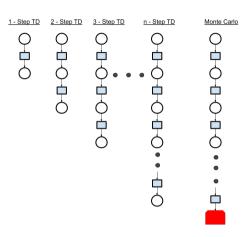
► n-step TD

n-step 1D
$$V(s_t) \leftarrow V(s_t) + \alpha_t [G_t^{(n)} - V(s_t)]$$



Why Multi-Step TD?





▶ Multi-step TD methods tend to have <u>better bias-variance tradeoff</u> compared to 1-step TD method

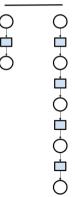
λ -Return



▶ What if we average some or all the *n*-step returns?

Example: Target could be

$$\frac{1}{2}G_t^{(2)} + \frac{1}{2}G_t^{(4)}$$



λ -Return



- ▶ Any set of returns can be averaged, even an infinite set, as long as the weights on the component returns are positive and sum to 1
- ▶ Choose $\lambda \in [0, 1]$, and define the λ -return at time t as

$$G_t^{\lambda} \stackrel{\text{def}}{=} (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

 \triangleright λ -return algorithm: use λ -return as the target

$$V(s_t) \leftarrow V(s_t) + \alpha_t [G_t^{\lambda} - V(s_t)]$$

▶ If the episode ends at T > t, define $G_t^{(n)}$ to be G_t for all n > T - t. Then

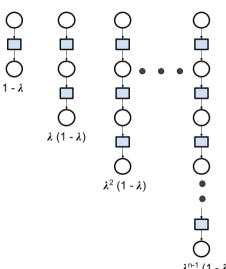
$$G_t^{\lambda} = (1 - \lambda) \sum_{n=1}^{T-t} \lambda^{n-1} G_t^{(n)} + \lambda^{T-t} G_t$$

 \rightarrow $\lambda = 0$ gives 1-step TD, $\lambda = 1$ gives MC update



λ -Return



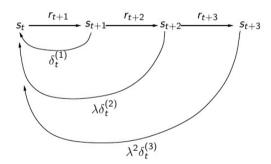




Forward View



$$G_t^{\lambda} - V(s_t) = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} [G_t^{(n)} - V(s_t)] = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} \delta_t^{(n)}$$

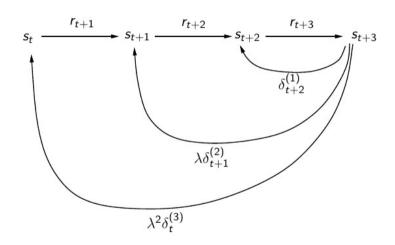


▶ Not suitable for online implementation;



A Possible Online Implementation





▶ Requires storing all rewards from the episode



A Rearrangement of *n*-step TD Errors

 $= \sum_{t+n-1}^{t+n-1} \gamma^{i-t} \delta_i^{(1)}$

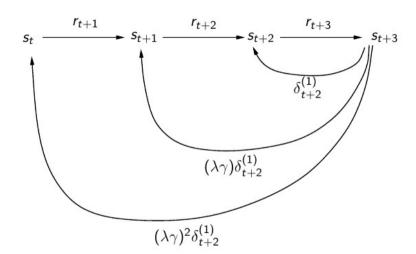


$$\delta_{t}^{(n)} \stackrel{\text{def}}{=} r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^{n} V(s_{t+n}) - V(s_{t})
= \gamma^{0} [r_{t+1} + \gamma V(s_{t+1}) - V(s_{t})]
+ \gamma^{1} [r_{t+2} + \gamma V(s_{t+2}) - V(s_{t+1})]
\vdots
+ \gamma^{n-1} [r_{t+n} + \gamma V(s_{t+n}) - V(s_{t+n-1})]$$

$$G_t^{\lambda} - V(s_t) = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} \delta_t^{(n)} = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} \sum_{i=t}^{t+n-1} \gamma^{i-t} \delta_i^{(1)}$$
$$= \sum_{n=1}^{\infty} (\lambda \gamma)^i \delta_{t+i}^{(1)}$$

Online Implementation





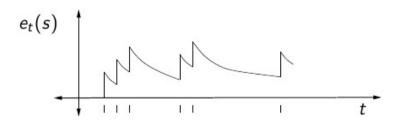
Eligibility Traces



▶ The eligibility trace of a state $s \in S$ at time t is defined recursively by

$$e_0(s) = 0$$

$$e_t(s) = \begin{cases} (\lambda \gamma)e_{t-1}(s), & s_t \neq s \\ (\lambda \gamma)e_{t-1}(s) + 1, & s_t = s \end{cases}$$



Algorithm : $TD(\lambda)$



Algorithm $TD(\lambda)$: Algorithm

- 1: Initialize e(s) = 0 for all s, V(s) arbitrarily
- 2: **for** For each episode **do**
- 3: Let s be a start state for episode k
- 4: for For each step of the episode do
- 5: Take action a recommended by policy π from state s
- Collect reward r and reach next state s'6:
- 7: Form the one-step TD error $\delta \leftarrow r + \gamma V(s') - V(s)$
- Increment eligibility trace of state $s, e(s) \leftarrow e(s) + 1$ 8:
- 9: for For all states $S \in \mathcal{S}$ do
- Update V(S): $V(S) \leftarrow V(S) + \alpha e(S)\delta$ 10:
- Update eligibility trace: $e(S) \leftarrow \lambda \gamma e(S)$ 11:
- end for 12:
- Move to next state: $s \leftarrow s'$ 13:
- 14: end for
- 15: end for



Certainity Equivalence Estimate



Certainity Equivalence Estimate



- Model based policy evaluation with model estimated from samples
- \blacktriangleright Given an experience quadruple (s, a, r, s'), we can estimate
 - \star Compute MLE for model using (s, a, s')

$$\hat{P}(s'|s,a) = \frac{1}{N(s,a)} \sum_{k=1}^{K} \sum_{t=1}^{L_k-1} \mathbb{1}(s_{k,t} = s, a_{k,t} = a, s_{k,t+1} = s')$$

★ Compute MLE estimate of the reward function

$$\hat{R}(s, a, s') = \frac{1}{N(s, a, s')} \sum_{k=1}^{K} \sum_{t=1}^{L_k - 1} \mathbb{1}(s_{k,t} = s, a_{k,t} = a, s_{k,t+1} = s') \ r_{t,k}$$

 \blacktriangleright Once MLE estimates of $\hat{P}(s'|s,a)$ and $\hat{R}(s,a,s')$ use a known policy evaluation method to compute a MLE based estimate for V^{π} .





Off Policy Learning



Off-Policy Learning



- ▶ Can we estimate the value V^{π} or Q^{π} of a target policy π ...
- ▶ While following a behavior policy μ ?
 - \star Trajectories or transitions are sampled from μ
 - \star Expected values have to be estimated w.r.t. π
- ▶ Possible benefits:
 - ★ Learn by observing other agents
 - ★ Re-use previous experience generated from earlier policies
 - ★ Learn about optimal policy while following exploratory policy
 - ★ Learn about multiple policies while following one policy

Importance Sampling: Review



Let P(x) be the <u>target</u> distribution and Q(x) be the <u>behaviour</u> distribution for some random variable x

$$\mathbb{E}_{X \sim P}[f(X)] = \sum_{i} P(x_i) f(x_i)$$

$$= \sum_{i} Q(x_i) \left[\frac{P(x_i)}{Q(x_i)} f(x_i) \right]$$

$$= \mathbb{E}_{X \sim Q} \left[\frac{P(X)}{Q(X)} f(X) \right]$$

 \blacktriangleright We have samples of X drawn from Q, but we wish to estimate expectation under P

$$\mathbb{E}_{X \sim P}[f(X)] \simeq \frac{1}{n} \sum_{i=1}^{n} \left[\frac{P(X_i)}{Q(X_i)} f(X_i) \right], \ X_i \sim Q$$

► Caveat: Q should not assign zero probability to any outcome that is assigned non-zero probability by P



Importance Sampling: Review



The ratio $\frac{P(x)}{Q(x)}$ is the importance sampling (IS) weight for x.

What about the variance of IS estimator $\widehat{\mu}_Q \approx \frac{1}{N} \sum_{x_i \in D} \left[\frac{P(x_i)}{Q(x_i)} f(x_i) \right]$?

$$\operatorname{var}_{Q}\left[\widehat{\mu}_{Q}\right] = \operatorname{var}_{Q}\left[\frac{P(x)}{Q(x)}f(x)\right]$$

$$= \left[\mathbb{E}_{Q}\left[\left(\frac{P(x)}{Q(x)}f(x)\right)^{2}\right] - \mathbb{E}_{Q}\left(\left[\frac{P(x)}{Q(x)}f(x)\right]\right)^{2}\right]$$

$$= \mathbb{E}_{P}\left[\left(\frac{P(x)}{Q(x)}f(x)^{2}\right)\right] - \mathbb{E}_{P}\left(f(x)\right)^{2}$$

If the likelihood ratio $\frac{P(x)}{O(x)}$ is large, the variance of the estimator explodes



Off-Policy TD using Importance Sampling



- ▶ Evaluate target policy π using TD targets $r_{t+1} + \gamma V(s_{t+1})$ generated from μ
- ▶ Weigh each TD target by the importance sampling factor $\pi(s_t, a_t)/\mu(s_t, a_t)$

$$V(s_t) \leftarrow V(s_t) + \alpha_t \left[\frac{\pi(s_t|a_t)}{\mu(s_t|a_t)} (r_{t+1} + \gamma V(s_{t+1})) - V(s_t) \right]$$

- \triangleright π may be deterministic and greedy, μ should be stochastic and exploratory
- ▶ The case $\mu = \pi$ is called *on-policy* learning
- \blacktriangleright On Policy Learning: Learn about policy π from experience sampled from π
- \triangleright Off Policy Learning: Learn about policy π from experience sampled from μ