

Model Free Control : TD Methods

Easwar Subramanian

TCS Innovation Labs, Hyderabad

Email : easwar.subramanian@tcs.com / cs5500.2020@iith.ac.in

October 04, 2021

- 1 Review
- 2 TD Control
- 3 Q-Learning
- 4 Summary and Closing Remarks

Review

- ▶ We use the principle behind policy iteration to do model free control as value iteration requires knowledge of model
 - ★ Policy evaluation : use MC or TD based model free prediction
 - ★ Policy improvement
- ▶ (Greedy) Policy improvement over V is also model based

$$\pi(s) = \arg \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V^\pi(s')]$$

- ▶ (Greedy) policy improvement over Q is model free

$$\pi(s) = \arg \max_a Q^\pi(s, a)$$

- ▶ For model-free policy improvement, we use Q^π , not V^π

Core Idea behind Model Free Control

- ▶ Initialize a policy π
- ▶ Repeat
 - ★ Policy Evaluation : Find Q^π
 - ★ Policy Improvement : Get an improved policy from evaluation of Q^π

- ▶ We now need to evaluate Q^π instead of V^π
- ▶ Recall that the state-action value function of a policy π is given by,

$$\begin{aligned} Q^\pi(s, a) &\stackrel{\text{def}}{=} \mathbb{E}_\pi(G_t | s_t = s, a_t = a) \\ &= \mathbb{E}_\pi \left(\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s, a_t = a \right) \\ &= \mathbb{E}_\pi(r_{t+1} + \gamma Q^\pi(s_{t+1}, a_{t+1}) | s_t = s, a_t = a) \end{aligned}$$


- ▶ We can use MC or TD methods to evaluate Q^π using samples

- ▶ To evaluate $Q^\pi(s, a)$ for some given state s and action a , repeat over several episodes
 - ★ The **first** time t that $s_t = s$ and $\pi(s) = a$ in the episode
 1. Increment counter for number of visits to s : $N(s, a) \leftarrow N(s, a) + 1$
 2. Increment running sum of total returns with return from current episode:
 $S(s, a) \leftarrow S(s, a) + G_t$
- ▶ Monte Carlo estimate of value function $Q(s, a) \leftarrow S(s, a)/N(s, a)$

The main drawback of this algorithm is

- ▶ Many state action pairs may never be visited
- ▶ If policy π is deterministic, things get even worse

Policy Iteration with Action Value Function



- Monte Carlo Policy Evaluation, $Q = Q^\pi$
- Greedy policy improvement, $\pi' = \arg \max_a Q^\pi(s, a)$

- ▶ Simplest idea for ensuring continual exploration
- ▶ All m actions are tried with non-zero probability every time
 - ★ With probability $1 - \varepsilon$, choose the greedy action
 - ★ With probability ε , choose an action uniformly at random

$$\begin{aligned}\pi(a|s) &= \frac{\varepsilon}{m} + 1 - \varepsilon, \text{ if } a = \arg \max_{a'} Q(s, a'), \\ &= \frac{\varepsilon}{m}, \text{ otherwise}\end{aligned}$$

ε -Greedy Policy Improvement

For any policy ε -greedy policy π , the ε -greedy policy π' w.r.t. Q^π is an improvement over π , that is, $V^{\pi'}(s) \geq V^\pi(s)$

$$\begin{aligned}Q^\pi(s, \pi'(s)) &= \sum_{a \in \mathcal{A}} \pi'(a|s) Q^\pi(s, a) \\&= \frac{\varepsilon}{m} \sum_{a \in \mathcal{A}} Q^\pi(s, a) + (1 - \varepsilon) \max_a Q^\pi(s, a) \\&= \frac{\varepsilon}{m} \sum_{a \in \mathcal{A}} Q^\pi(s, a) + (1 - \varepsilon) \frac{1 - \varepsilon}{1 - \varepsilon} \max_a Q^\pi(s, a) \\&= \frac{\varepsilon}{m} \sum_{a \in \mathcal{A}} Q^\pi(s, a) + (1 - \varepsilon) \sum_a \frac{\pi(a|s) - \frac{\varepsilon}{m}}{1 - \varepsilon} \max_a Q^\pi(s, a) \\&\geq \frac{\varepsilon}{m} \sum_{a \in \mathcal{A}} Q^\pi(s, a) + (1 - \varepsilon) \sum_a \frac{\pi(a|s) - \frac{\varepsilon}{m}}{1 - \varepsilon} Q^\pi(s, a) \\&= \sum_{a \in \mathcal{A}} \pi(a|s) Q^\pi(s, a) = V^\pi(s)\end{aligned}\tag{1}$$

Therefore, $V^{\pi'}(s) \geq V^\pi(s)$ from the policy improvement theorem

Definition


Greedy in the Limit with Infinite Exploration

- ▶ All state-action pairs are visited infinitely often
- ▶ The policy converges to a purely greedy policy

$$\lim_{k \rightarrow \infty} \pi_k(a|s) = \mathbf{1}_{a=\arg \max_{a'} Q_k(s,a)}$$

- ▶ ε -greedy is GLIE if ε decays to 0 asymptotically, for example,

$$\varepsilon_k = \frac{1}{k}$$



Every episode

- Monte Carlo Policy Evaluation $Q \approx Q^\pi$
- Policy improvement using ϵ - greedy with ϵ decay

Algorithm Monte Carlo Control : GLIE

- 1: Initialize $Q(s,a) = 0$, set $\epsilon = 1$;
- 2: Create an ϵ -greedy initial policy π_1 ;
- 3: **for** $k = 1, 2, \dots, K$ **do**
- 4: Sample a trajectory from policy π_k
- 5: **for** For each state action (s_t, a_t) pair in the trajectory **do**
- 6: Compute the total discounted return G_t starting from (s_t, a_t)
- 7:

$$N(s_t, a_t) = N(s_t, a_t) + 1$$

- 8:
- 9: **end for**
- 10: Set $\epsilon \leftarrow \frac{1}{k}$ and perform the policy improvement step as

$$\pi_{k+1} = \epsilon\text{-greedy}(\pi_k)$$

- 11: **end for**
-

TD Control

- ▶ Natural idea : Use TD instead of MC in policy iteration framework
- ▶ Apply TD to evaluate $Q(s, a)$ in the evaluation step
- ▶ Use ε -greedy policy improvement in the update step

- State-action value function of a policy π :

$$\begin{aligned} Q^\pi(s, a) &\stackrel{\text{def}}{=} \mathbb{E}_\pi(G_t | s_t = s, a_t = a) \\ &= \mathbb{E}_\pi(r_{t+1} + \gamma Q^\pi(s_{t+1}, a_{t+1}) | s_t = s, a_t = a) \end{aligned}$$

- Iterative DP policy evaluation:

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} \mathcal{P}_{ss'}^a \left[\mathcal{R}_{ss'}^a + \gamma \sum_{a'} (\pi(s', a') Q_k(s', a')) \right]$$
$$Q_k \rightarrow Q^\pi$$

- TD approximation: Given the transition $(s_t, a_t, r_{t+1}, s_{t+1})$, sample $a' \sim \pi(s_{t+1}, \cdot)$, and update


$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha_t [r_{t+1} + \gamma Q(s_{t+1}, a') - Q(s_t, a_t)]$$

- ▶ TD approximation: Given the transition $(s_t, a_t, r_{t+1}, s_{t+1})$, sample $a' \sim \pi(s_{t+1}, \cdot)$, and perform the following update

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha_t[r_{t+1} + \gamma Q(s_{t+1}, a') - Q(s_t, a_t)]$$

- ▶ On-policy version (SARSA): $a_t \sim \pi(s_t, \cdot)$
- ▶ Off-policy version: $a_t \sim \mu(s_t, \cdot)$;
 - ★ Need to multiply the term inside square brackets with suitable importance sampling factor

- ▶ On Policy and off-policy version converges to Q^π
 - ★ Convergence takes place under similar conditions as TD methods for V^π
 - ▶ State and action spaces are finite
 - ▶ All state-action pairs are visited infinitely often
 - ▶ Robbins-Monroe condition: $\sum_t \alpha_t = \infty, \sum_t \alpha_t^2 < \infty$



Along every episode, we interleave one step of policy evaluation followed ϵ -greedy policy improvement

- Policy is always ε -greedy with ε decay
- Given a trajectory segment $(s_t, a_t, r_{t+1}, s_{t+1}, a_{t+1})$ generated by the ε -greedy policy, update

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha_t[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

Algorithm SARSA

- 1: Initialize $Q(s, a)$ arbitrarily, with Q at terminal states set to zero
- 2: **for** Repeat for each episode **do**
- 3: Initialize s , choose action a at s using ϵ -greedy over Q
- 4: **for** Repeat for each step in the episode **do**
- 5: Take action a , observe reward r and next state s'
- 6: Choose action a' for state s' using ϵ -greedy over Q
- 7:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)], s \leftarrow s', a \leftarrow a'$$

- 8: **end for**
 - 9: **end for**
-

Q-Learning

- Optimal Q function:

$$Q_*(s, a) \stackrel{\text{def}}{=} \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- Bellman optimality equation:

$$Q_*(s, a) = \mathbb{E} \left[r_{t+1} + \gamma \max_{a'} Q_*(s_{t+1}, a') \mid s_t = s, a_t = a \right]$$

- Iterative DP approximation

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} \mathcal{P}_{ss'}^a \left[\mathcal{R}_{ss'}^a + \gamma \max_{a'} Q_k(s', a') \right]$$

$$Q_k \rightarrow Q_*$$

- Policy is always ε -greedy with ε decay
- Given a trajectory segment $(s_t, a_t, r_{t+1}, s_{t+1})$ generated by the ε -greedy policy, update

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha_t [r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t)]$$

Algorithm Q-Learning

- 1: Initialize $Q(s, a)$ arbitrarily, with Q at terminal states set to zero
- 2: **for** Repeat for each episode **do**
- 3: Initialize s , choose action a at s using ε -greedy over Q
- 4: **for** Repeat for each step in the episode **do**
- 5: Take action a , observe reward r and next state s'
- 6: Choose target to update $Q(s, a)$ by being greedy at s' as shown below
- 7:


$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)], s \leftarrow s'$$

- 8: **end for**
- 9: **end for**

SARSA and Q-Learning : Backup diagram

- Q-learning is an off-policy algorithm
 - ★ Target policy is greedy w.r.t to $Q(s, a)$,
 - ★ Behaviour policy is ϵ -greedy w.r.t to $Q(s, a)$

Backup Diagrams for SARSA and Q-Learning



Summary and Closing Remarks

- ▶ MC-based evaluation of V^π (also possible for Q^π)
- ▶ TD-based approximate evaluation of V^π, Q^π
 - ★ 1-step TD, n -step TD, TD(λ), SARSA, Q -learning
 - ★ Convergence guarantees under infinite exploration, and Robbins-Monroe condition
- ▶ TD-based control
 - ★ On-policy control with SARSA (also possible: n -step SARSA, SARSA(λ))
 - ★ Off-policy control with Q -learning
 - ★ Based on optimistic policy iteration, and GLIE

- ▶ TD methods have several advantages over MC methods
 - ★ Lower variance
 - ★ Online
 - ★ Partial sequences

Schematic View of MC and TD Algorithms






Figure: MC Algorithm and TD Algorithms



- ▶ Tic-Tac-Toe : States : Board positions and moves are actions
- ▶ A conventional action-value function ($Q(s, a)$) would map or learn about the two state action pairs on the top row separately
- ▶ An afterstate value function would immediately evaluate both equally
- ▶ Any learning about the position-move pair on the left would immediately transfer to the pair on the right

Pavlov's Dog and Temporal Difference




Pavlov's Dog

$$V(s) \leftarrow V(s) + \alpha[r + \gamma V(s') - V(s)]$$

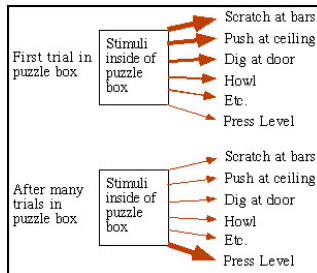
Figure Source:

<https://www.age-of-the-sage.org/psychology/pavlov.html>

Thondrike's Cat and Exploration



Thondrike's cat



Actions by cat

ϵ -greedy strategy helps to explore !!

All methods discussed under model free methods are in the tabular setting

- ▶ Next: richer ways to represent value functions
- ▶ Needed for very large (or continuous) state spaces
- ▶ What if the action space is large (or continuous)?

Over to Deep RL !!